



**BITS Pilani**  
Pilani | Dubai | Goa | Hyderabad

## DEEP NEURAL NETWORK MODULE # 3 : LINEAR NEURAL NETWORKS FOR REGRESSION

---

Seetha Parameswaran  
BITS Pilani WILP

The instructor is gratefully acknowledging  
the authors who made their course  
materials freely available online.

This deck is prepared by Seetha Parameswaran.



# TABLE OF CONTENTS

- 1 REGRESSION
- 2 LINEAR REGRESSION
- 3 THE LINEAR MODEL: SINGLE NEURON
- 4 WORKED EXAMPLE
- 5 EVALUATION
- 6 IMPLEMENTATION TIPS
- 7 SUMMARY AND EXTENSIONS



# WHAT IS REGRESSION?

## DEFINITION

Regression is a supervised learning task that predicts a continuous-valued output based on input features.

## Key Characteristics:

- Input: Feature vector  $\mathbf{x} \in \mathbb{R}^d$
- Output: Real-valued target  $y \in \mathbb{R}$
- Goal: Learn a function  $f : \mathbb{R}^d \rightarrow \mathbb{R}$  such that  $y \approx f(\mathbf{x})$

# EXAMPLES OF REGRESSION PROBLEMS

## ① House Price Prediction

- ▶ Input: Size, location, number of rooms
- ▶ Output: Price in dollars

## ② Temperature Forecasting

- ▶ Input: Historical weather data, season, humidity
- ▶ Output: Temperature in degrees

## ③ Stock Price Prediction

- ▶ Input: Historical prices, trading volume, market indicators
- ▶ Output: Future stock price

# TABLE OF CONTENTS

1 REGRESSION

2 LINEAR REGRESSION

3 THE LINEAR MODEL: SINGLE NEURON

4 WORKED EXAMPLE

5 EVALUATION

6 IMPLEMENTATION TIPS

7 SUMMARY AND EXTENSIONS



# WHAT IS LINEAR REGRESSION?

## DEFINITION


Linear regression assumes the output is a linear combination of the input features.

Mathematical Form:

$$y = w_0 + w_1x_1 + w_2x_2 + \cdots + w_dx_d \quad (1)$$

where:

- $\mathbf{x} = [x_1, x_2, \dots, x_d]^T$  are the input features
- $\mathbf{w} = [w_1, w_2, \dots, w_d]^T$  are the weights
- $w_0$  is the bias term



# EXAMPLES OF LINEAR REGRESSION

## ① Salary Prediction

- ▶  $\text{Salary} = w_0 + w_1 \times \text{Years of Experience}$

## ② Crop Yield Prediction

- ▶  $\text{Yield} = w_0 + w_1 \times \text{Rainfall} + w_2 \times \text{Fertilizer}$


## ③ Energy Consumption

- ▶  $\text{Energy} = w_0 + w_1 \times \text{Temperature} + w_2 \times \text{Hour} + w_3 \times \text{Day}$

The relationship between inputs and output is assumed to be **linear**.

# VISUAL EXAMPLE: 2D LINEAR REGRESSION

2D Linear Regression: Scatter Plot with Fitted Line



# TABLE OF CONTENTS

- 1 REGRESSION
- 2 LINEAR REGRESSION
- 3 THE LINEAR MODEL: SINGLE NEURON
- 4 WORKED EXAMPLE
- 5 EVALUATION
- 6 IMPLEMENTATION TIPS
- 7 SUMMARY AND EXTENSIONS



# THE FOUR COMPONENTS

A complete machine learning system consists of:

- ① Data:  $d$ -dimensional input vectors and target outputs
- ② Model: Single neuron implementing linear function
- ③ Objective Function: Measures prediction error
- ④ Learning Algorithm: Gradient descent to optimize weights


# MATRIX FORM OF DATA

For efficient computation, we organize data into matrices:

**Input Data** = Augmented Features:

Add bias term:  $\tilde{\mathbf{x}}^{(i)} = [1, x_1^{(i)}, \dots, x_d^{(i)}]^T$

Each row = one training example


Design Matrix:  $\mathbf{X} \in \mathbb{R}^{N \times (d+1)}$    $\mathbf{X} = \begin{bmatrix} 1 & x_1^{(1)} & \dots & x_d^{(1)} \\ 1 & x_1^{(2)} & \dots & x_d^{(2)} \\ \vdots & \vdots & \ddots & \vdots \\ 1 & x_1^{(N)} & \dots & x_d^{(N)} \end{bmatrix}$

Target Vector:  $\mathbf{y} \in \mathbb{R}^N$   $\mathbf{y} = [y^{(1)} \ y^{(2)} \ \dots \ y^{(N)}]^T$

Weight Vector:  $\mathbf{w} \in \mathbb{R}^{d+1}$   $\mathbf{w} = [w_0 \ w_1 \ w_2 \ \dots \ w_d]^T$

# LINEAR MODEL AS A SINGLE NEURON

Model linear regression as a single artificial neuron:



- Input:  $d$ -dimensional feature vector  $\mathbf{x} \in \mathbb{R}^d$
- Output: Single predicted value  $\hat{y} \in \mathbb{R}$

# LINEAR MODEL PREDICTION

Model: A single neuron computes a weighted sum and applied identity activation function.

Prediction for input  $\mathbf{x}$ :

$$\hat{y} = f(\mathbf{w}^T \mathbf{x}) = f(w_0 + w_1 x_1 + w_2 x_2 + \cdots + w_d x_d) \quad (2)$$

where:

- $\mathbf{w} = [w_0, w_1, \dots, w_d]^T \in \mathbb{R}^{d+1}$  are the model parameters
- $\mathbf{x} = [1, x_1, \dots, x_d]^T \in \mathbb{R}^{d+1}$  is the input
- $f(\cdot)$  is the identity activation  $f(z) = z$

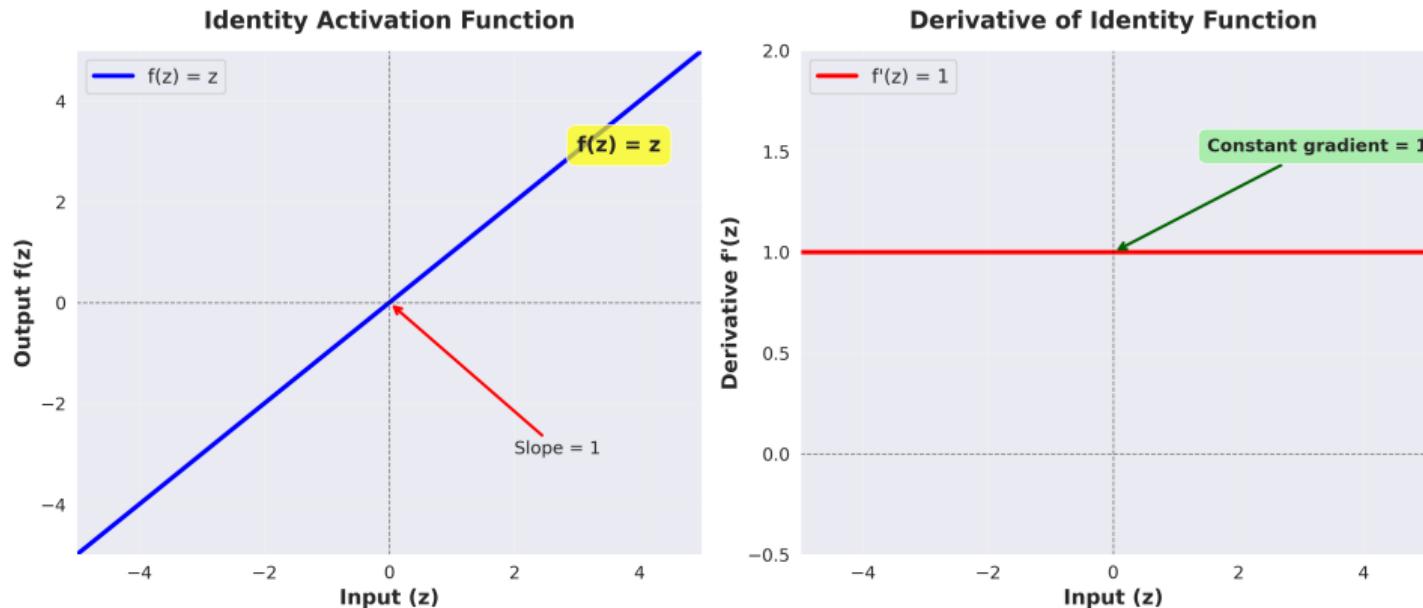
Vectorized form: for all predictions

$$\hat{\mathbf{y}} = \mathbf{X}\mathbf{w} \quad (3)$$

# IDENTITY ACTIVATION FUNCTION

## Identity Activation Function

$$f(z) = z \quad (4)$$



# IDENTITY ACTIVATION FUNCTION

## Properties of Identity Function

- Output: Continuous,  $(-\infty, \infty)$
- Differentiable everywhere:  $f'(z) = 1$
- Allows output to take any real value
- Gradient flows unchanged through the neuron
- Perfect for predicting continuous targets

## Why identity for regression?

- Target values are continuous (prices, temperatures, etc.)
- Need to predict any real number, not just binary { 0,1 }
- Differentiability enables gradient-based optimization
- No information loss from input to output

# OBJECTIVE FUNCTION: SQUARED ERROR LOSS

We need to measure how well our model fits the data.

Loss for single example:

$$\ell(\mathbf{w}; \mathbf{x}^{(i)}, y^{(i)}) = \frac{1}{2} (\hat{y}^{(i)} - y^{(i)})^2 = \frac{1}{2} (\mathbf{w}^T \mathbf{x}^{(i)} - y^{(i)})^2 \quad (5)$$

Total Loss (Mean Squared Error):

$$J(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N \frac{1}{2} (\mathbf{w}^T \mathbf{x}^{(i)} - y^{(i)})^2 \quad (6)$$

Goal: Find  $\mathbf{w}^*$  that minimizes  $J(\mathbf{w})$

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} J(\mathbf{w}) \quad (7)$$

# WHY SQUARED ERROR?

Advantages of squared error loss:

- **Differentiable:** Smooth everywhere, easy to optimize
- **Convex:** Single global minimum (for linear models)
- **Penalizes large errors:** Quadratic growth
- **Statistical interpretation:** Maximum likelihood under Gaussian noise

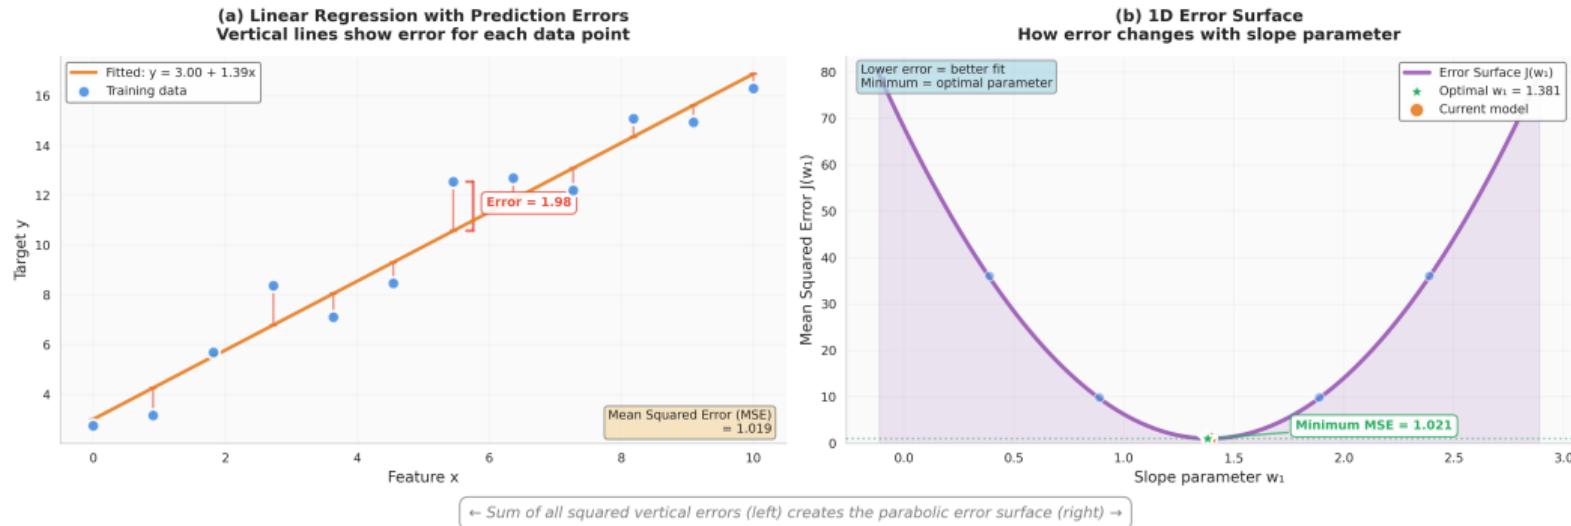
Vector Form:

$$J(\mathbf{w}) = \frac{1}{2N} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2 = \frac{1}{2N} (\mathbf{X}\mathbf{w} - \mathbf{y})^T (\mathbf{X}\mathbf{w} - \mathbf{y}) \quad (8)$$

where  $\mathbf{X} \in \mathbb{R}^{N \times (d+1)}$  is the design matrix and  $\mathbf{y} \in \mathbb{R}^N$  is the target vector.

# VISUALIZING THE ERROR SURFACE

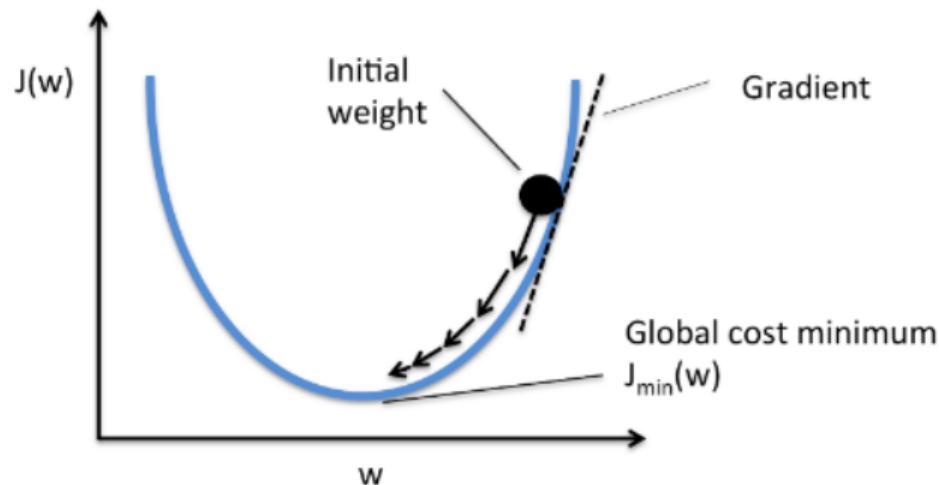
## Understanding Error Surface: From Individual Errors to Loss Function



The error surface is **convex** with a single global minimum.

# GRADIENT DESCENT ALGORITHM

Idea: Iteratively update weights in the direction of steepest descent



where  $\eta > 0$  is the learning rate (step size).

Intuition: Move downhill on the error surface to find the minimum.

# BATCH GRADIENT DESCENT ALGORITHM

---

## Algorithm 1: Gradient Descent for Linear Regression

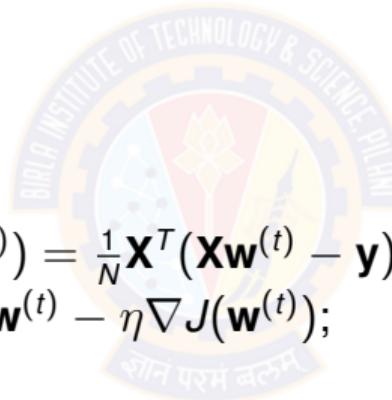
---

**Input:** Dataset  $\mathcal{D} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^N$ , learning rate  $\eta$ , tolerance  $\epsilon$

**Output:** Learned weights  $\mathbf{w}$

```
1 Initialize  $\mathbf{w}^{(0)} = \mathbf{0}$  (or random);
2  $t \leftarrow 0$ ;
3 while not converged do
4     Compute gradient:  $\nabla J(\mathbf{w}^{(t)}) = \frac{1}{N} \mathbf{X}^T (\mathbf{X}\mathbf{w}^{(t)} - \mathbf{y})$ ;
5     Update weights:  $\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \eta \nabla J(\mathbf{w}^{(t)})$ ;
6     if  $\|\nabla J(\mathbf{w}^{(t+1)})\| < \epsilon$  then
7         break // Converged
8      $t \leftarrow t + 1$ ;
9 return  $\mathbf{w}^{(t+1)}$ ;
```

---



# COMPUTING THE GRADIENT

Gradient of squared error loss:

$$\nabla J(\mathbf{w}) = \frac{\partial J}{\partial \mathbf{w}} = \frac{1}{N} \sum_{i=1}^N (\mathbf{w}^T \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)} \quad (9)$$

Matrix Form:

$$\nabla J(\mathbf{w}) = \frac{1}{N} \mathbf{X}^T (\mathbf{X}\mathbf{w} - \mathbf{y}) \quad (10)$$

where  $\mathbf{X}^T \in \mathbb{R}^{(d+1) \times N}$ ,  $\mathbf{X}\mathbf{w} - \mathbf{y} \in \mathbb{R}^N$ , so  $\nabla J \in \mathbb{R}^{d+1}$

# GRADIENT DESCENT UPDATE RULE

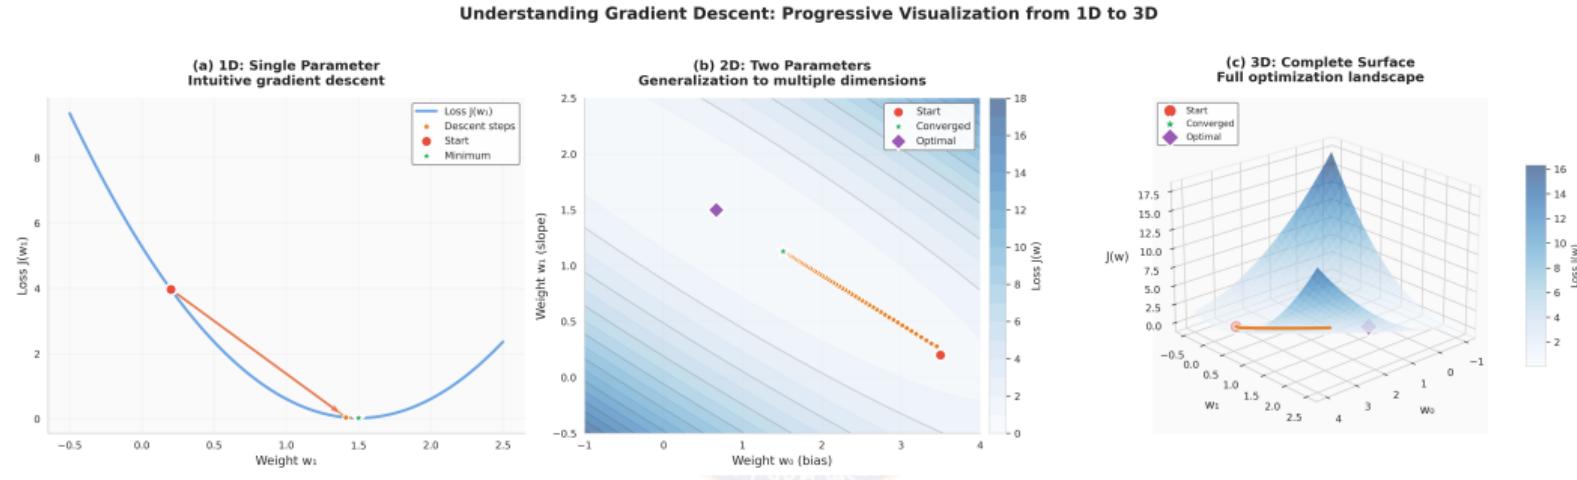
Update equation:

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \frac{\eta}{N} \mathbf{X}^T (\mathbf{X}\mathbf{w}^{(t)} - \mathbf{y}) \quad (11)$$

Key Parameters:

- Learning rate  $\eta$ : Controls step size
  - ▶ Too large: May overshoot minimum
  - ▶ Too small: Slow convergence
- Stopping criterion:  $\|\nabla J(\mathbf{w})\| < \epsilon$  or max iterations

# VISUALIZING GRADIENT DESCENT



Each step moves perpendicular to contour lines toward the minimum.

# TABLE OF CONTENTS

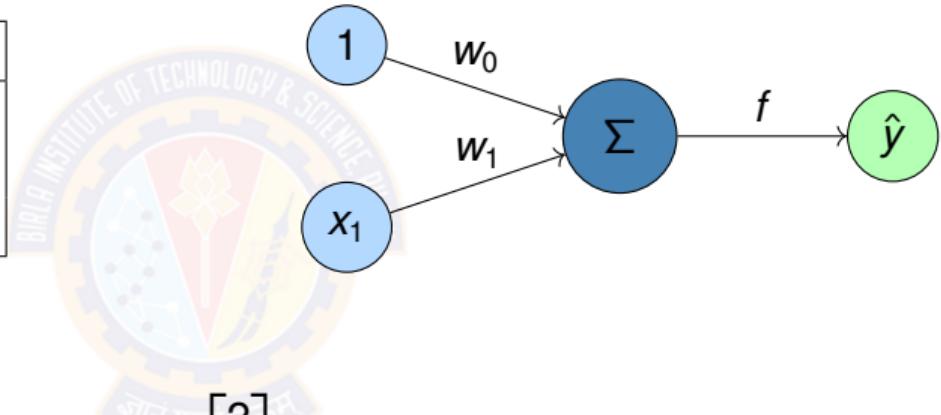
- 1 REGRESSION
- 2 LINEAR REGRESSION
- 3 THE LINEAR MODEL: SINGLE NEURON
- 4 WORKED EXAMPLE
- 5 EVALUATION
- 6 IMPLEMENTATION TIPS
- 7 SUMMARY AND EXTENSIONS



# NUMERICAL EXAMPLE: SETUP

Tiny Dataset: Predict house price from size

Size ( $x_1$ )	Price ( $y$ )
1	2
2	4
3	5



Matrix Formulation:

$$\mathbf{X} = \begin{bmatrix} 1 & 1 \\ 1 & 2 \\ 1 & 3 \end{bmatrix} \in \mathbb{R}^{3 \times 2}, \quad \mathbf{y} = \begin{bmatrix} 2 \\ 4 \\ 5 \end{bmatrix} \in \mathbb{R}^3, \quad \mathbf{w} = \begin{bmatrix} w_0 \\ w_1 \end{bmatrix} \in \mathbb{R}^2$$

# NUMERICAL EXAMPLE: INITIALIZATION

Hyperparameters:  $\eta = 0.1$   $N = 3$

Initialize weights:  $\mathbf{w}^{(0)} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$

Vectorized Prediction:  $\hat{\mathbf{y}} = \mathbf{X}\mathbf{w}$

Initial predictions:  $\hat{\mathbf{y}}^{(0)} = \mathbf{X}\mathbf{w}^{(0)} = \begin{bmatrix} 1 & 1 \\ 1 & 2 \\ 1 & 3 \end{bmatrix} \begin{bmatrix} 0 \\ 0 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}$

Loss Function:  $J(\mathbf{w}) = \frac{1}{2N} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2$

Initial loss:  $J(\mathbf{w}^{(0)}) = \frac{1}{6} \left\| \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix} - \begin{bmatrix} 2 \\ 4 \\ 5 \end{bmatrix} \right\|^2 = \frac{1}{6}(4 + 16 + 25) = 7.5$

# NUMERICAL EXAMPLE: COMPUTING GRADIENT

Error vector:  $\mathbf{e}^{(0)} = \hat{\mathbf{y}}^{(0)} - \mathbf{y} = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix} - \begin{bmatrix} 2 \\ 4 \\ 5 \end{bmatrix} = \begin{bmatrix} -2 \\ -4 \\ -5 \end{bmatrix}$

Gradient computation:  $\nabla J(\mathbf{w}^{(0)}) = \frac{1}{N} \mathbf{X}^T \mathbf{e}^{(0)}$

$$\nabla J(\mathbf{w}^{(0)}) = \frac{1}{3} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 2 & 3 \end{bmatrix} \begin{bmatrix} -2 \\ -4 \\ -5 \end{bmatrix}$$

$$= \frac{1}{3} \begin{bmatrix} -2 - 4 - 5 \\ -2 - 8 - 15 \end{bmatrix} = \begin{bmatrix} -3.67 \\ -8.33 \end{bmatrix}$$

# NUMERICAL EXAMPLE: WEIGHT UPDATE

Gradient descent update:  $\mathbf{w}^{(1)} = \mathbf{w}^{(0)} - \eta \nabla J(\mathbf{w}^{(0)})$

Computing new weights:  $\mathbf{w}^{(1)} = \begin{bmatrix} 0 \\ 0 \end{bmatrix} - 0.1 \begin{bmatrix} -3.67 \\ -8.33 \end{bmatrix} = \begin{bmatrix} 0.367 \\ 0.833 \end{bmatrix}$

New predictions:  $\hat{\mathbf{y}}^{(1)} = \mathbf{X}\mathbf{w}^{(1)}$

$$\hat{\mathbf{y}}^{(1)} = \begin{bmatrix} 1 & 1 \\ 1 & 2 \\ 1 & 3 \end{bmatrix} \begin{bmatrix} 0.367 \\ 0.833 \end{bmatrix} = \begin{bmatrix} 1.20 \\ 2.03 \\ 2.87 \end{bmatrix}$$

# NUMERICAL EXAMPLE: NEW LOSS

New error vector:  $\mathbf{e}^{(1)} = \hat{\mathbf{y}}^{(1)} - \mathbf{y} = \begin{bmatrix} 1.20 \\ 2.03 \\ 2.87 \end{bmatrix} - \begin{bmatrix} 2 \\ 4 \\ 5 \end{bmatrix} = \begin{bmatrix} -0.80 \\ -1.97 \\ -2.13 \end{bmatrix}$


Loss after one iteration:  $J(\mathbf{w}^{(1)}) = \frac{1}{6} \|\mathbf{e}^{(1)}\|^2 = \frac{1}{6} (0.64 + 3.88 + 4.54) \approx 1.51$

Progress:

- Initial loss:  $J(\mathbf{w}^{(0)}) = 7.5$
- After 1 iteration:  $J(\mathbf{w}^{(1)}) = 1.51$
- 80% reduction! ✓**

Continue until  $\|\nabla J(\mathbf{w})\| < \epsilon$  or max iterations reached

# COMPUTATIONAL GRAPH FOR ONE GRADIENT



# TABLE OF CONTENTS

- 1 REGRESSION
- 2 LINEAR REGRESSION
- 3 THE LINEAR MODEL: SINGLE NEURON
- 4 WORKED EXAMPLE
- 5 EVALUATION
- 6 IMPLEMENTATION TIPS
- 7 SUMMARY AND EXTENSIONS



# MODEL EVALUATION: TRAINING VS TEST ERROR

**Key Principle:** Evaluate on unseen data

Data Split:

- **Training set:** Used to learn weights  $\mathbf{w}$  (typically 90-99%)
- **Test set:** Used to evaluate generalization (typically 1-10%)

$$\text{Training Error: } J_{\text{train}} = \frac{1}{N_{\text{train}}} \sum_{i \in \text{train}} (\hat{y}^{(i)} - y^{(i)})^2$$

$$\text{Test Error: } J_{\text{test}} = \frac{1}{N_{\text{test}}} \sum_{i \in \text{test}} (\hat{y}^{(i)} - y^{(i)})^2$$

**Goal:** Minimize test error (generalization performance)

# EVALUATION METRICS

1. Mean Squared Error (MSE):

$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (\hat{y}^{(i)} - y^{(i)})^2 \quad (12)$$

2. Root Mean Squared Error (RMSE):

$$\text{RMSE} = \sqrt{\text{MSE}} = \sqrt{\frac{1}{N} \sum_{i=1}^N (\hat{y}^{(i)} - y^{(i)})^2} \quad (13)$$

3. Mean Absolute Error (MAE):  
**(Less sensitive to outliers than MSE)**

$$\text{MAE} = \frac{1}{N} \sum_{i=1}^N |y^{(i)} - \hat{y}^{(i)}| \quad (14)$$

# $R^2$ SCORE (COEFFICIENT OF DETERMINATION)

Definition: Proportion of variance explained by the model

$$R^2 = 1 - \frac{\sum_{i=1}^N (y^{(i)} - \hat{y}^{(i)})^2}{\sum_{i=1}^N (y^{(i)} - \bar{y})^2} = 1 - \frac{SS_{\text{residual}}}{SS_{\text{total}}} \quad (15)$$

where  $\bar{y} = \frac{1}{N} \sum_{i=1}^N y^{(i)}$  is the mean.

Interpretation:

- $R^2 = 1$ : Perfect fit (all variance explained)
- $R^2 = 0$ : Model no better than predicting mean
- $R^2 < 0$ : Model worse than baseline (possible on test set)
- Typical:  $0.7 < R^2 < 0.9$  indicates good fit

# TABLE OF CONTENTS

- 1 REGRESSION
- 2 LINEAR REGRESSION
- 3 THE LINEAR MODEL: SINGLE NEURON
- 4 WORKED EXAMPLE
- 5 EVALUATION
- 6 IMPLEMENTATION TIPS
- 7 SUMMARY AND EXTENSIONS



# TIPS FOR IMPLEMENTATION

## 1. Choosing Learning Rate $\eta$ :

- Start with  $\eta \in \{0.001, 0.01, 0.1, 1.0\}$
- Plot loss vs iterations for each
- If loss increases or becomes NaN  $\rightarrow$  decrease  $\eta$
- If convergence too slow  $\rightarrow$  increase  $\eta$
- Use learning rate schedules for fine-tuning

## 2. When to Stop Training:

- Gradient magnitude:  $\|\nabla J\| < \epsilon$  (e.g.,  $\epsilon = 10^{-4}$ )
- Loss change:  $|J^{(t)} - J^{(t-1)}| < \epsilon$
- Maximum iterations reached (e.g.,  $T = 1000$  or  $10,000$ )
- Validation loss starts increasing (early stopping for overfitting)

# MORE PRACTICAL TIPS

## 3. Weight Initialization:

- Zero initialization:  $\mathbf{w}^{(0)} = \mathbf{0}$  (works well for linear regression)
- Avoid large initial values (can cause divergence or overflow)

## 4. Feature Engineering:

- Domain knowledge: Create meaningful derived features (e.g., price per sq ft)
- Feature selection: Remove irrelevant/redundant features


## 5. Feature Scaling and Normalization:

- Features with different scales can slow convergence
- 1. Min-Max Scaling: Scale to  $[0, 1]$   $x'_j = \frac{x_j - \min(x_j)}{\max(x_j) - \min(x_j)}$
- 2. Standardization (Z-score): Zero mean, unit variance  $x'_j = \frac{x_j - \mu_j}{\sigma_j}$

# DEBUGGING CHECKLIST

## Common Issues and Solutions:

- Loss is NaN/Inf:
  - ▶ Learning rate too large
  - ▶ Numerical overflow (check feature scales)
  - ▶ Solution: Decrease  $\eta$ , scale features
- Training works but test error high:
  - ▶ Overfitting (too many features relative to data)
  - ▶ Solution: Regularization, more data, or simpler model
- Loss oscillating:
  - ▶ Learning rate too large
  - ▶ Solution: Decrease  $\eta$  or use learning rate decay
- Loss not decreasing:
  - ▶ Learning rate too small
  - ▶ Bug in gradient computation (check signs!)
  - ▶ Wrong sign in update rule (should be minus)
  - ▶ Solution: Increase  $\eta$ , verify code




# TABLE OF CONTENTS

- 1 REGRESSION
- 2 LINEAR REGRESSION
- 3 THE LINEAR MODEL: SINGLE NEURON
- 4 WORKED EXAMPLE
- 5 EVALUATION
- 6 IMPLEMENTATION TIPS
- 7 SUMMARY AND EXTENSIONS



# SUMMARY: PUTTING IT ALL TOGETHER



*This framework extends to more complex models: neural networks, deep learning!*


# KEY TAKEAWAYS

- Regression predicts continuous outputs from input features.
- Linear regression assumes a linear relationship:  $y = \mathbf{w}^T \mathbf{x} + w_0$
- A single neuron with identity activation implements linear regression.
- Squared error loss provides a convex objective function.
- Gradient descent iteratively optimizes weights:  $\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla J$
- Feature scaling is critical for fast convergence.
- Evaluation on test set ensures generalization.

# LINEAR REGRESSION

**Key insight:** Linear relationship between inputs and output, optimized by minimizing squared errors.

- **Problem:** Predict continuous values:  $y \in \mathbb{R}$
- **Model:** Single neuron with identity activation
- **Output:** Predicted value  $\hat{y} = \mathbf{w}^T \mathbf{x} \in (-\infty, \infty)$
- **Loss:** Mean Squared Error (MSE)
- **Training:** Batch Gradient Descent (GD)



# REFERENCES

- Zhang, A., Lipton, Z. C., Li, M., & Smola, A. J. (2023). *Dive into Deep Learning*. Cambridge University Press. Chapter 3 (Linear Neural Networks for Regression). <https://d2l.ai/>
- Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep Learning*. MIT Press. Chapter 5 (Machine Learning Basics).  
<https://www.deeplearningbook.org/>
- Bishop, C. M. (2006). *Pattern Recognition and Machine Learning*. Springer. Chapter 3 (Linear Models for Regression).
- Murphy, K. P. (2022). *Probabilistic Machine Learning: An Introduction*. MIT Press. Chapter 11 (Linear Regression).  
<https://probml.github.io/pml-book/>

**Thank You!**