



*AIML CLZG516
ML System Optimization
Session 2*

Parallel Programming Models:

- Pipe-lined, Data-Parallel, Task-Parallel, and Request-Parallel
- Speedup

Add two numbers:-

a = memory [100].

b = memory [101]

c = a + b

memory [102] = c.

(index memory loc.)

RAM

LOAD R1, [100].

LOAD R2, [101]

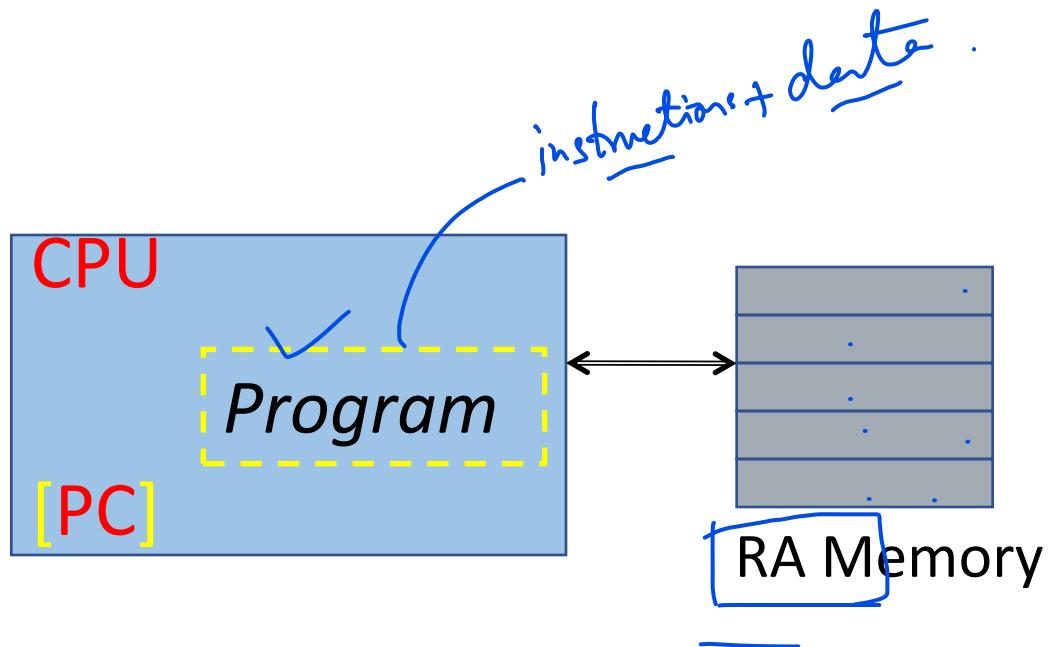
ADD R3, R1, R2

STORE R3, [102]

HALT.

Algorithm Design - Sequential

- Generic Machine Model
 - Random Access Machine Model



Typical Instructions

- Arithmetic / logic operations,
- Load / Store, and
- Jump / Branch

PC: Program Counter
(tracks the next instruction to be executed)

RAM: Random Access Memory
(cost of access is uniform across locations)

Executing an Instruction

- Different stages of Instruction Execution:

<u>Fetch Instruction</u>	Move the next instruction (tracked by PC) to a register
<u>Decode Instruction</u>	Identify Operator and (data) addresses
<u>Load (data)</u>	Move data into register (if needed)
<u>Execute</u>	Perform the operation
<u>Store (result)</u>	Move the resulting data to memory

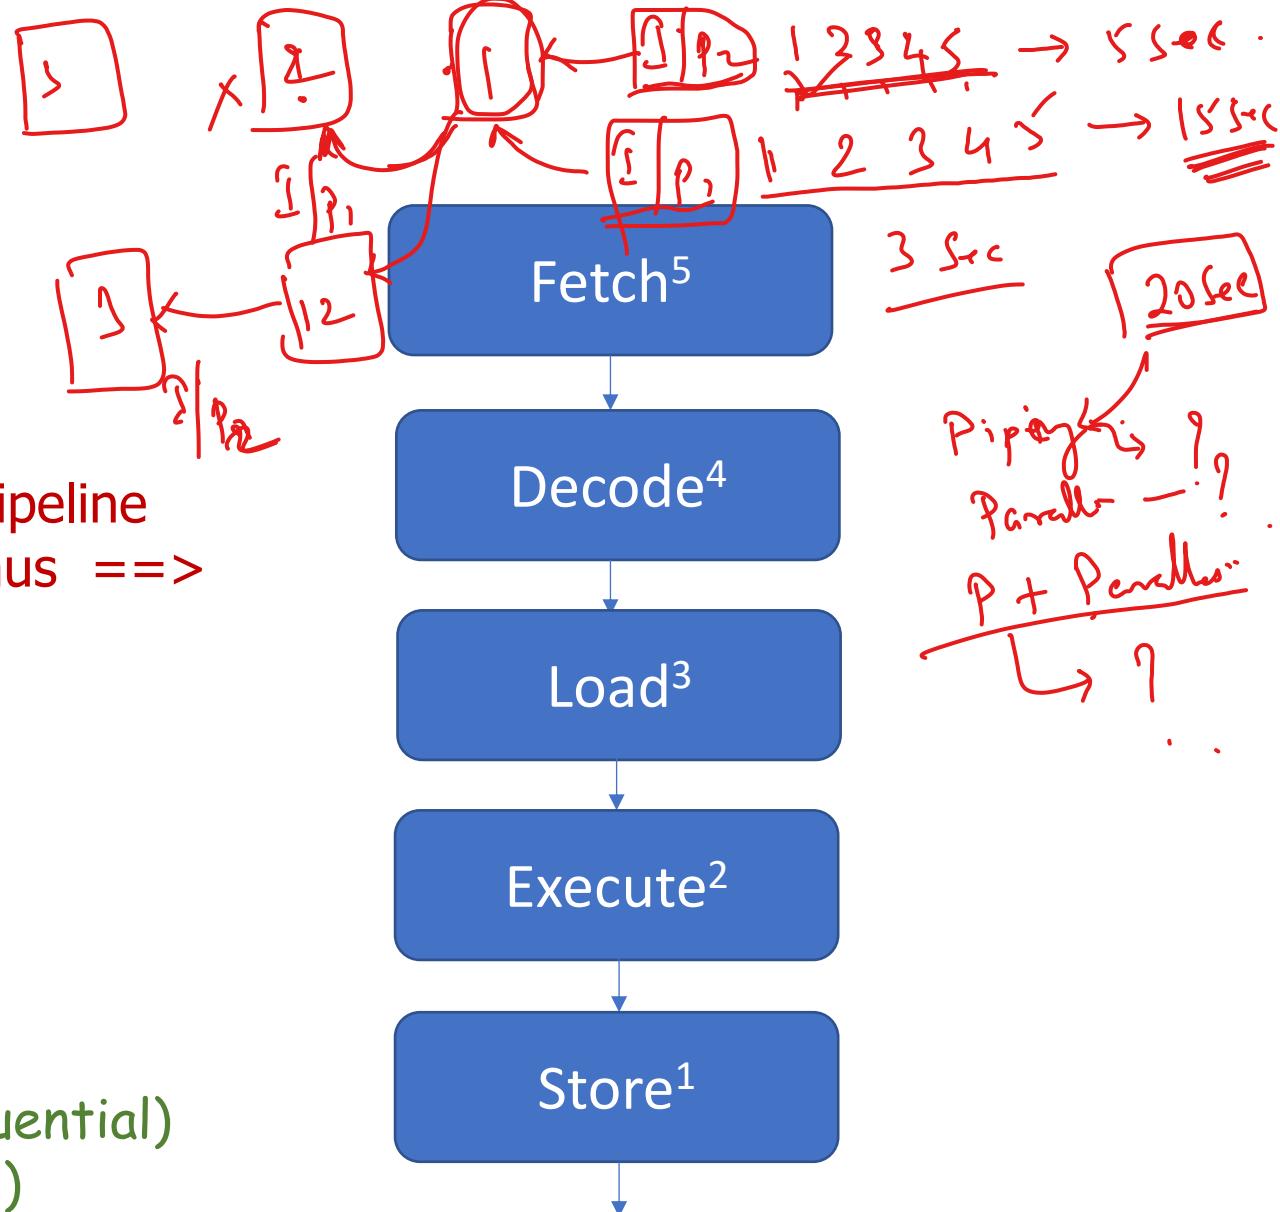
If separate circuitry is designed for each stage-
so that the stage take the same amount of time -
then a sequence of instructions can be executed in a pipeline

Instruction Pipeline

Given a sequence of instructions of the form:

- I1
- I2
- I3
- I4
- I5
- ...

execution in a pipeline would appear thus ==>



If each stage takes 1 clock cycle, then throughput has increased:

- from 1 instruction per 5 cycles (sequential)
- to 1 instruction per 1 cycle (pipelined)

Q: What about Turn-around-Time aka response time?

Modern Processors

- Modern processors (since Intel Pentium circa 1991)
 - typically include a pipeline that is several stages (>5) deep
- Throughput in processors is measured in CPI (or Cycles Per Instruction):
 - For an ideal pipeline design: CPI is 1
 - In practice, it may be more (Why?)
 - but on an average it is kept close to 1

Pipeline Throughput

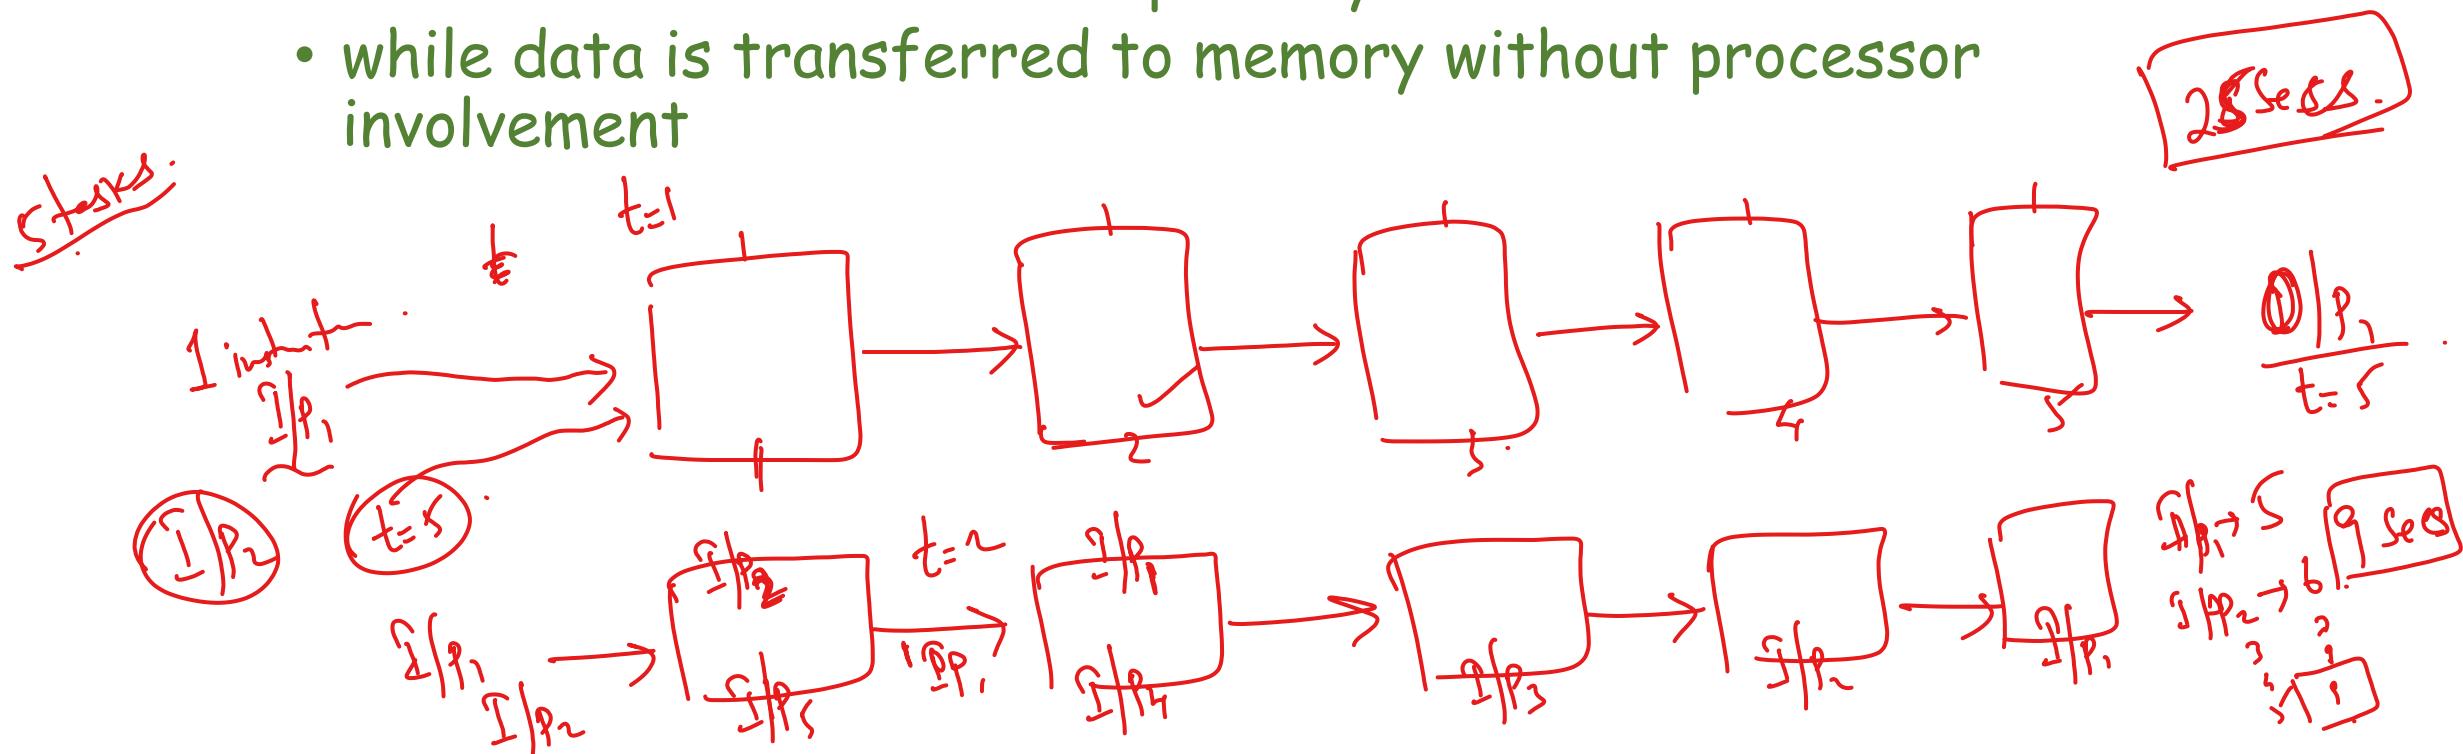
- Speedup (i.e. throughput increase) is k for a k -stage pipeline
- Factors that may slow down the pipeline:
 - Some stage(s) take more time than others
 - Q: What is the impact on CPI if one stage takes 10% extra time?
 - Memory access takes more time (i.e., LOAD and STORE)
 - Modern processor pipelines are designed
 - such that all stages take almost the same time (except for LOAD and STORE)

Pipeline Throughput and Memory Access [2]

- Memory access is slower compared to Processor speed:
 - Typical processor clock cycles
 - e.g. 2 to 3 GHz
 - i.e., in-processor operation may take only 0.33 to 0.5ns
 - Access from Memory (DRAM) will take around
 - 50 to 200ns
 - Access from Cache (SRAM) - if available - may take
 - 5 to 10ns.
- Modern architectures use multiple levels of caching and other techniques to keep the access time low.
- Compiler and processor collaborate to keep the frequency of memory access operations low.

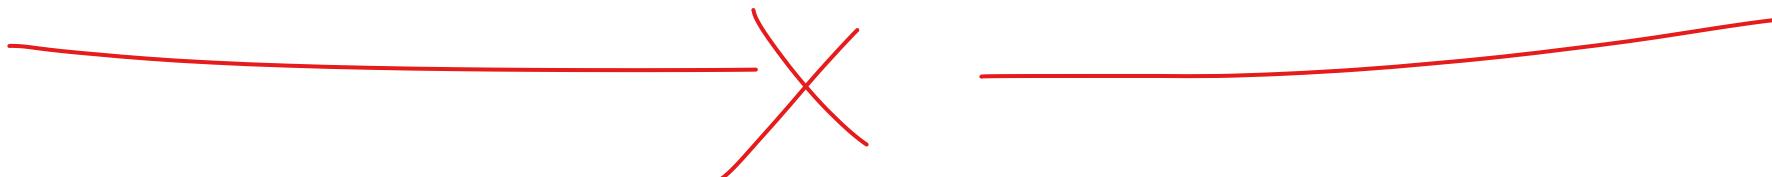
Pipeline Throughput and Memory Access

- STORE operations may be executed asynchronously:
 - i.e. processor does not wait for data to be stored in memory
 - Store buffers (i.e., buffer registers inside the processor) are used to store the data temporarily
 - while data is transferred to memory without processor involvement



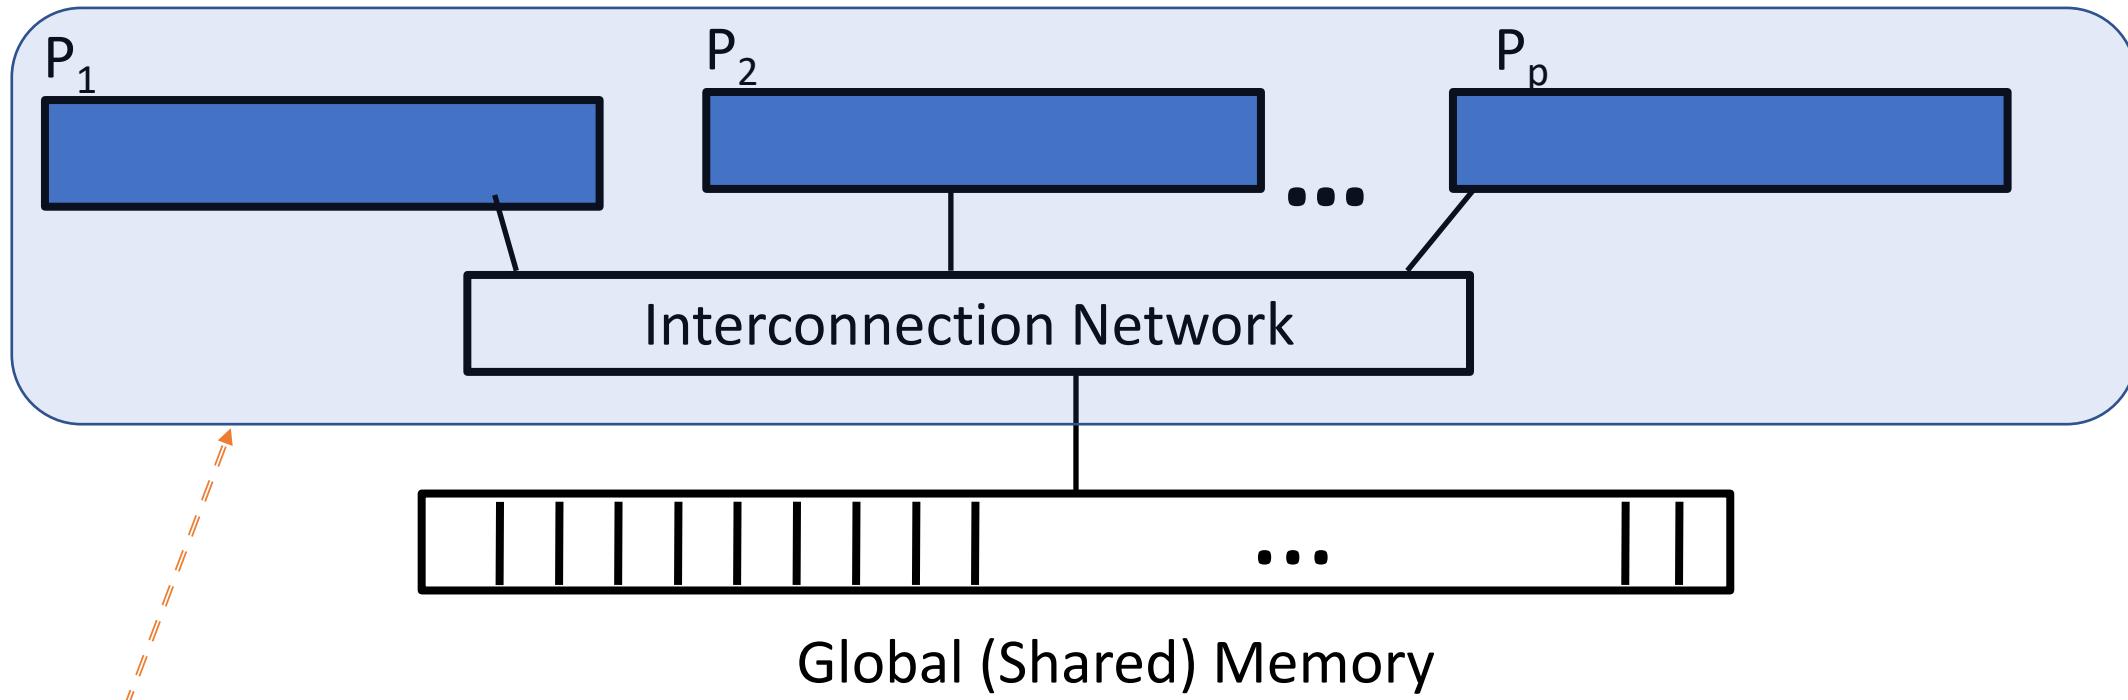
Software Pipelining

- The idea of a pipeline can be extended to Software Design:
 - ✓ • Break a long task into multiple stages
 - so that the stages take (roughly) the same amount of time.
 - If there is a stream of data to be processed by the data,
 - then the stream can be fed to the pipeline for improved throughput.
 - We will revisit this later!



Algorithm Design - Parallel: Shared Memory Model

Target environment:



e.g. a multi-core chip

Multi-threaded Programming:
each thread runs on a separate core

Typical Instructions

- Arithmetic / logic operations,
- Load / Store, and
- Jump / Branch

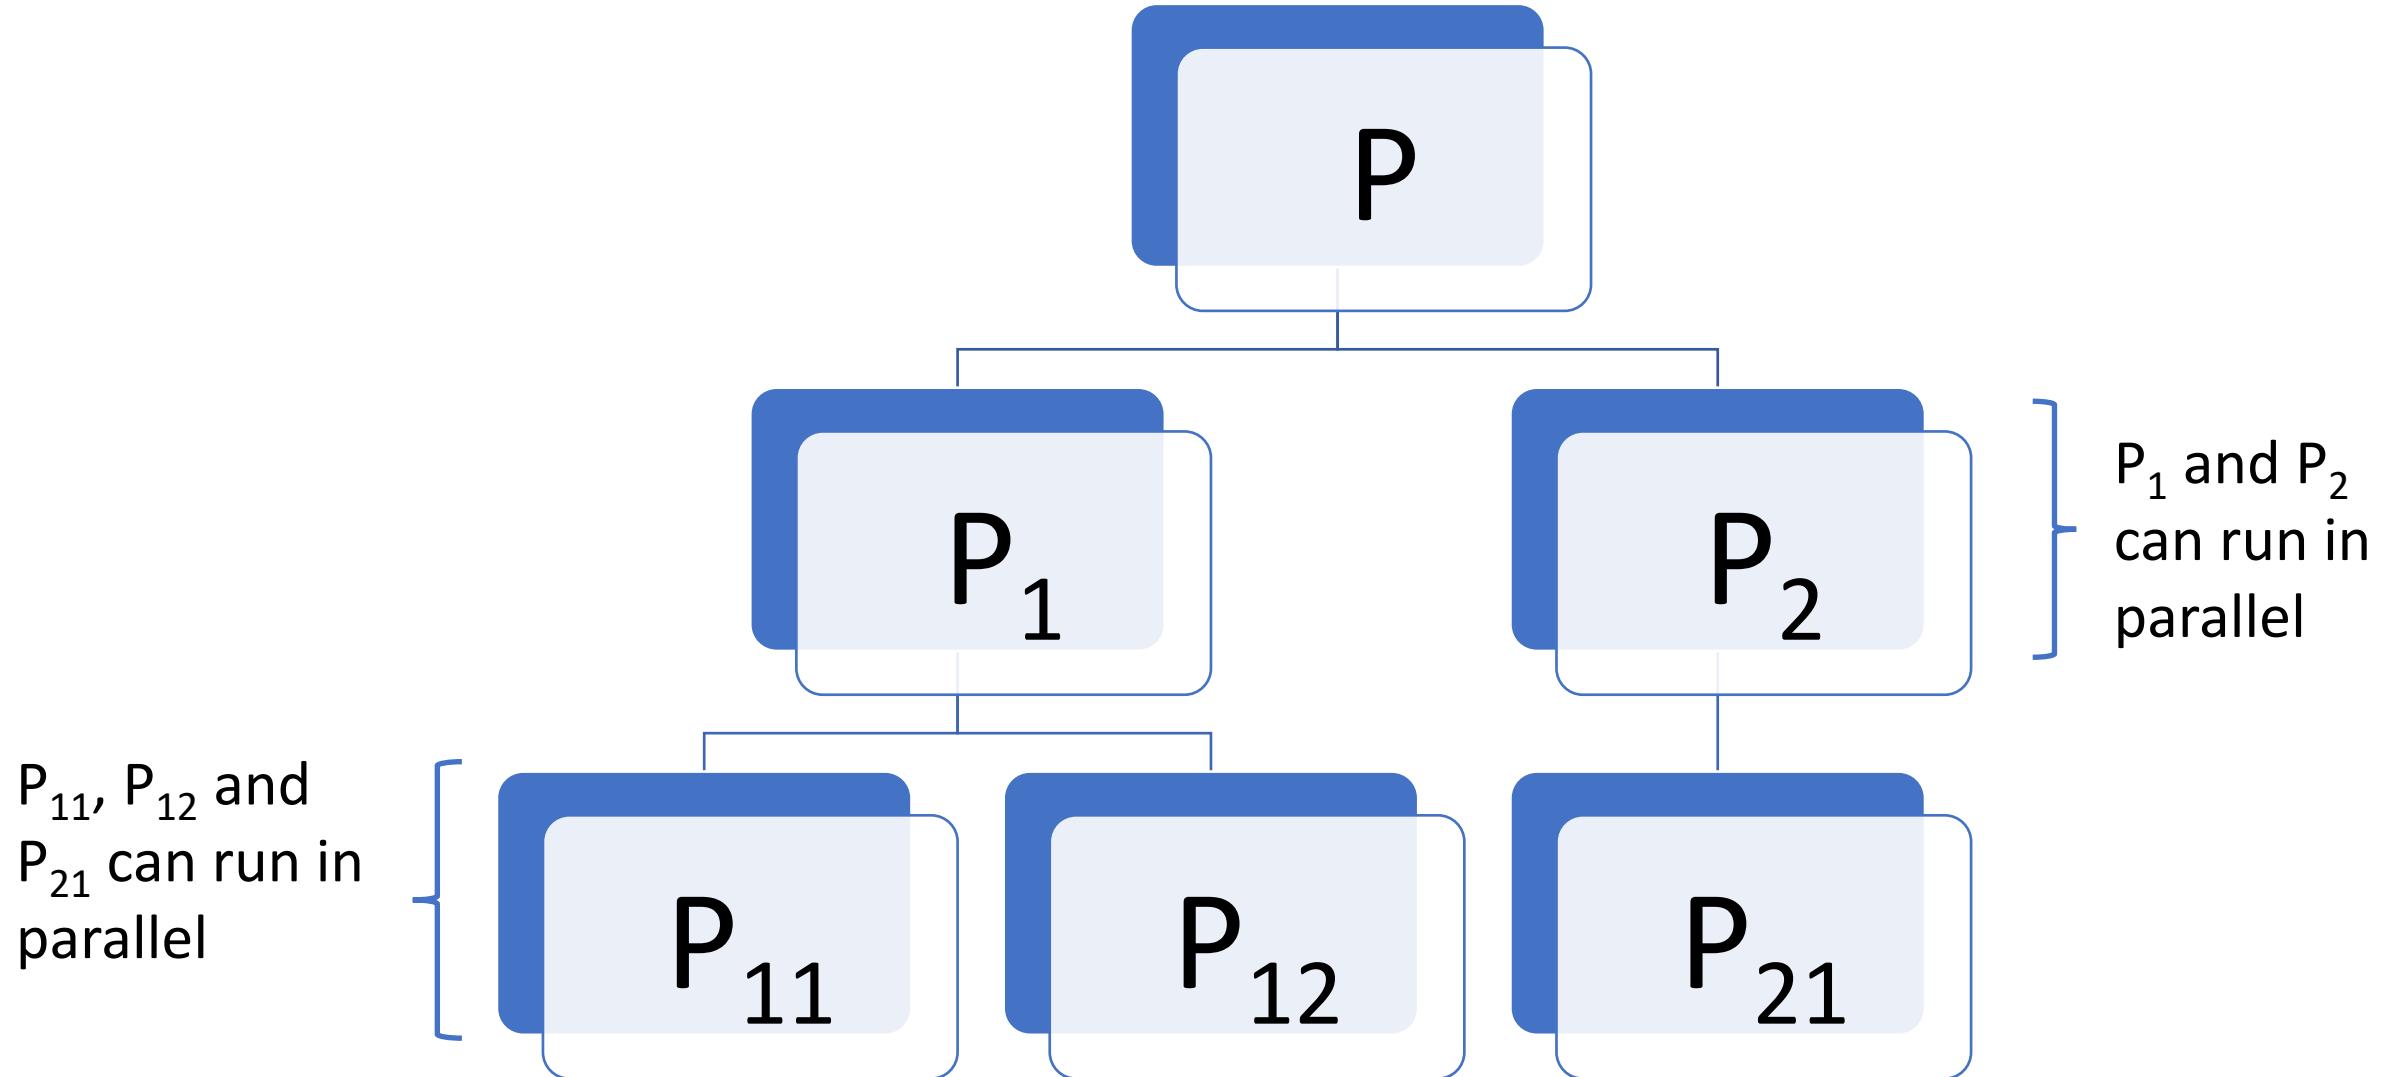
Algorithm Design

- Top-Down Design (Top Down Decomposition)
 1. Divide the problem into sub-problems.
 2. Find solutions for sub-problems
 3. Combine the sub-solutions.
- How do we find solutions for sub problems?
 - Apply top-down design recursively (i.e. divide each sub-problem further)
 - Q: When do we stop dividing?
 - A: When we reach "atomic" problems.
 - Atomic problems have known solutions
- Does any decomposition work?
 - Divide (the problem) only if you know how to combine (the solutions)

Top Down Design - Parallel

- Does any decomposition work?
 - Divide (the problem) only if you know how to combine (the solutions)
 - But also:
 - Mapping sub-problems to processors
 - Where is the combination done?
 - Number of sub-problems?
 - Processor utilization is the key!
 - i.e. More the number of processors more the number of sub-problems!

Top-Down Design - Parallel



Example: Search a key k in a list L_s of size N

Data: Assume $L_s[0..N-1]$ is stored in shared memory

t is N/p

for processor P_j from $j = 0$ to $p-1$

do $res_j = \text{search}(k, L_s[j*t..(j+1)*t-1])$

for processor P_0 :

do $res = \text{TRUE}$;

for $j = 0$ to $p-1$ do $res = res \text{ AND } res_j$



This is an example of data parallel programming!

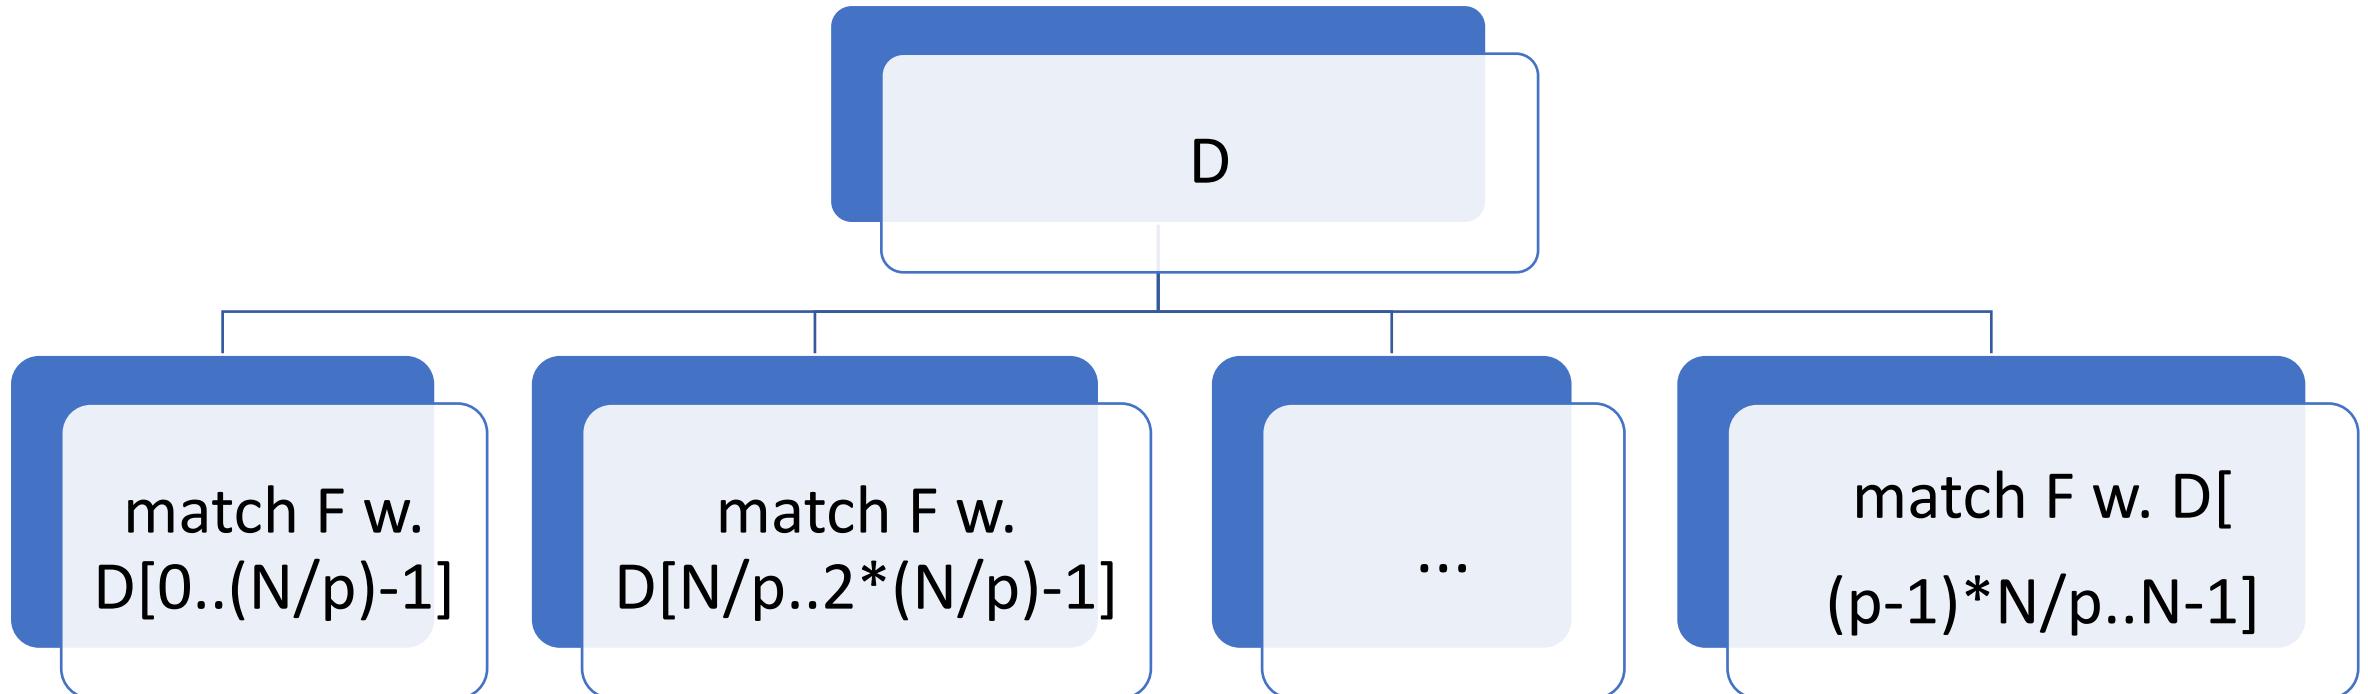
✓ Data Parallel

- Data Parallel execution (or computation):
 - The same task executes independently (i.e., in parallel) on different data
 - i.e., divide given data into (roughly) equal-sized subsets
 - and the same task is replicated and run on different processors - one for each subset.
- Note that we are assuming a shared memory model
 - i.e., all processors can access the (global) shared memory
 - Dividing data may simply become setting (boundary) markers!
- This may result in memory contention:
 - i.e., performance may not scale (with number of processors)

Data Parallel Execution - Example

Fingerprint Matching:

- Match a given print F with a database D of prints available



Data Parallel Execution - Exercise

- Vector Product $A \cdot B$ for two vectors - each of length N
- $\sum_{j=1 \text{ to } N} A[j] * B[j]$
- N processors:
 - For each processor $j=1 \text{ to } N$ do: $A[j] * B[j]$
- How to do the addition? Can it be done in data-parallel fashion?
- p processors: Change the code!

✓ SPMD

- Data-Parallel execution is also referred to as
 - Single Program Multiple Data (SPMD) programming (because a single program i.e. the same program) is executed on all processors
- This model Data-Parallel or SPMD is preferred where feasible
 - because of ease of programming and efficiency.
- In the parallel programming world, efficiency is measured as *speedup*:
 - i.e., the ratio of time taken by a parallel algorithm to time taken by a sequential algorithm

Speedup

- Speedup (in running time) of a given algorithm A running on p processors is defined as:
 - $\text{Speedup}(p) = (\text{Time taken by } A \text{ on 1 processor}) / (\text{Time taken by } A \text{ on } p \text{ processors})$
- All parts of a program may not run independently or in parallel:
 - Memory contention
 - Data (structure) contention
 - Mutually exclusive access (e.g. update operations or transactions) of shared data
 - Data dependency (result of a task must be input to another)

✓ Speedup - Amdahl's Law

- Assume that a fraction f of a task is not parallelizable (e.g., due to constraints seen in the last slide)
- $\text{Speedup}(p) = 1/(f + (1-f)/p)$
 - i.e., the parallelizable fraction $(1-f)$ of the program has been sped-up by a factor of p , the number of processors
 - But the other part takes the same (fraction of) time f
- By definition, $f=0$ in data-parallel execution or an SPMD program:
 - and $\text{speedup}(p) = p$
- When $\text{speedup}(p)$ is proportional to p , we say that the algorithm is scalable.