



BITS Pilani
WILP

AIML CLZG516
ML System Optimization
Murali Parameswaran





*AIML CLZG516
ML System Optimization
Session 8*

Distributed ML - Models and Platforms

- Implementation Issues
- The Parameter Server Model
- Stochastic Gradient Descent

Decision Trees

- Approach:
 - Construct a tree where each node denotes a binary decision
 - Nodes in the tree correspond to features and the order of features is chosen based on the notion of information gain (IG)
 - Information gain is the entropy
 - entropy of the whole set
 - minus the entropy when a particular feature is chosen

Decision Tree Construction

- Algorithm ID3 (input dataset S)
 - If all examples have the same label
 - Return a leaf with that label
 - Else if there are no features left to test
 - Return a leaf with the most common label
 - Else choose the feature F that maximizes IG of dataset S as the next node
 - Add a branch from the node for each possible value f in F
 - For each branch:
 - Calculate S_f by removing F from the set of features
 - $ID3(S_f)$

Parallel/Distributed Tree Construction?

- When you branch assign each branch (corresponding to one value of a feature)
 - To a different task (task parallelism)
 - At each level : number of parallel tasks = number of possible values of a feature

ML problem and Error

- Given input dataset - a vector of size n ,
 - Each training example x_i of d features (or dimensions) is associated
 - with a label y_i and
 - model parameters (likely corresponding to the features)
 - The problem is to predict y corresponding to an unseen example x
- Training error
 - Difference between the predicted y the actual label y' for an x

Model complexity

- Relation between model size (number of parameters) and data size (for training):
 - If there is too little data,
 - then a highly detailed model may overfit
 - If the model is too small,
 - then it may fail to capture relevant attributes
- Regularization addresses this issue

ML as regularized error minimization

- Training an ML model is minimizing the function F :
 - $F(w) = \sum_i L(x_i, y_i, w) + \Omega(w)$
 - where w denotes the set of parameters and
 - L is the loss function (i.e. prediction error) and
 - Ω is the regularizer that penalizes the model for complexity

Distributed ML


- Complexity:
 - Training Data size: from 1TB to 1PB
 - Model Size: 10^9 to 10^{12} parameters
- Examples:
 - Online Recommender System
 - millions of user profiles
 - Ad click predictor
 - each training example is a feature vector of high dimensions

Distributed ML - System Requirements

- In a distributed system,
 - the training data is partitioned among multiple nodes
 - and the nodes together learn the parameter vector w .
- The algorithm operates iteratively:
 - In each iteration,
 - every node independently uses its own training data to
 - Compute the changes to be made to w in order to move closer to an optimal value
 - Each node computes changes to w based only on its local data,
 - a central place is needed to aggregate these changes

Distributed Systems - Blackboard architecture


- Blackboard architecture is a pattern for distributed computation
 - where multiple nodes have to combine results
 - computed locally, in parallel - see processes P_i below



Regularized Error Minimization : A Distributed Architecture

In each iteration:

- Each worker node W_i pulls (current) parameters w from the server, computes $F(w)$, and posts it back.
- Server updates and minimizes $F(w)$




Distributed Systems and Failures

- Individual Nodes may fail frequently in distributed systems:
 - This is particularly so in commodity clusters
- Rate of node failure increases with the size of the system (i.e., more nodes and more processes)
 - Cloud data centers are made out of commodity clusters!
- Distributed Systems have to function (i.e. be available) despite node failures
 - This is referred to as fault tolerance and is achieved via
 - redundancy and failure recovery

Scalable and Dependable (*reliable and available*) architecture for ML

Each Worker Group includes a task scheduler.

Server Group must be fault tolerant!



Scalable and Dependable (*reliable and available*) architecture for ML

- This architecture is referred to as the **parameter server** model:
 - Different servers may handle different problems
 - Multiple servers may work on the same problem to improve performance
 - This will require additional combination/minimization processes and servers.
 - Multiple servers may work on the same problem for redundancy.

Parameter Server Model

- This model was popular for a few years
 - Google built an internal platform named DistBelief based on this model
 - DistBelief was optimized primarily for large clusters of multi-core nodes
 - GPU clusters were enabled later
 - Later, Google's TensorFlow provided programming flexibilities not available in DistBelief:
 - Adding new layers
 - Adding new ML training workflows
 - Optimizing or tuning ML algorithms

TensorFlow

- GPU acceleration has become a common tool for ML algorithms.
 - Building and testing on GPU workstations before scaling it to a GPU cluster is a common scenario as well.
- TensorFlow provides a unified programming interface and a common runtime on all these hardware platforms
 - while also supporting heterogeneous accelerators.
- e.g. Google's TPUs are special purpose accelerators for ML
 - that enable increased performance-per-watt compared to other state-of-the-art hardware.
- TensorFlow supports a common device abstraction for heterogeneous accelerators.

Gradient Descent

- One approach to error minimization is known as Gradient Descent:
 - Use the slope (i.e., gradient) of the loss function to update the parameters.
- This is particularly useful in Neural Networks in the back-propagation phase

Gradient Descent

- The gradient descent approach to minimize the error requires the following update to the parameters
 - $w = w - \eta * g(L, D, w)$
 - where g is the gradient function, L the loss function, and D , the dataset.
 - η denotes the learning rate - controls the amount by which the parameters are updated.
 - The updates are done iteratively.

Gradient Descent

```
for i = 1 to num_iter :
```

1. grad = eval_gradient (loss_function , D , w)
2. w = w - learning_rate * grad

- This is Batch GD!
 - i.e., update is done after all the points in the dataset are considered.
- Batch Gradient Descent is slow to converge, particularly if same data (or similar) data is repeated within the dataset.:

Stochastic Gradient Descent (SGD)

```
1. for i = 1 to num_iter :  
    1.1 shuffle ( data )  
    1.2 for example in data :  
        1.2.1 grad = eval_gradient ( loss_function , example , w )  
        1.2.2. w = w - learning_rate * grad
```

This may converge faster but is completely sequential i.e., not easy to parallelize!

Mini-batch SGD

```
for i = 1 to num_iter:  
    1. shuffle ( data )  
    2. for batch in get_batches (data , batch_size): → Parallelize /  
        1. grad = eval_gradient ( loss_function , batch , w )  
        2. w = w - learning_rate * grad  
            distribute
```

- Batches in the inner loop can be executed independently (locally)!
 - If necessary, batches can be obtained and stored locally at the start.
- Update has to be done in the parameter server