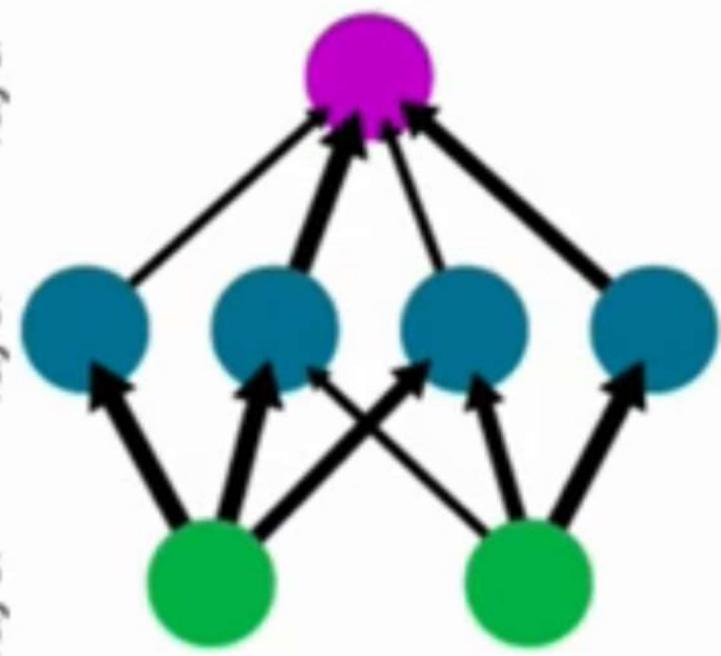


What is Backpropagation?

Backpropagation is a supervised learning algorithm, for training a neural network.

A simple neural network

input layer
hidden layer
output layer



Backpropagation

The backpropagation algorithm consists of two parts:

- forward pass
- backward pass

In the forward pass, the expect output corresponding to the given inputs are evaluated

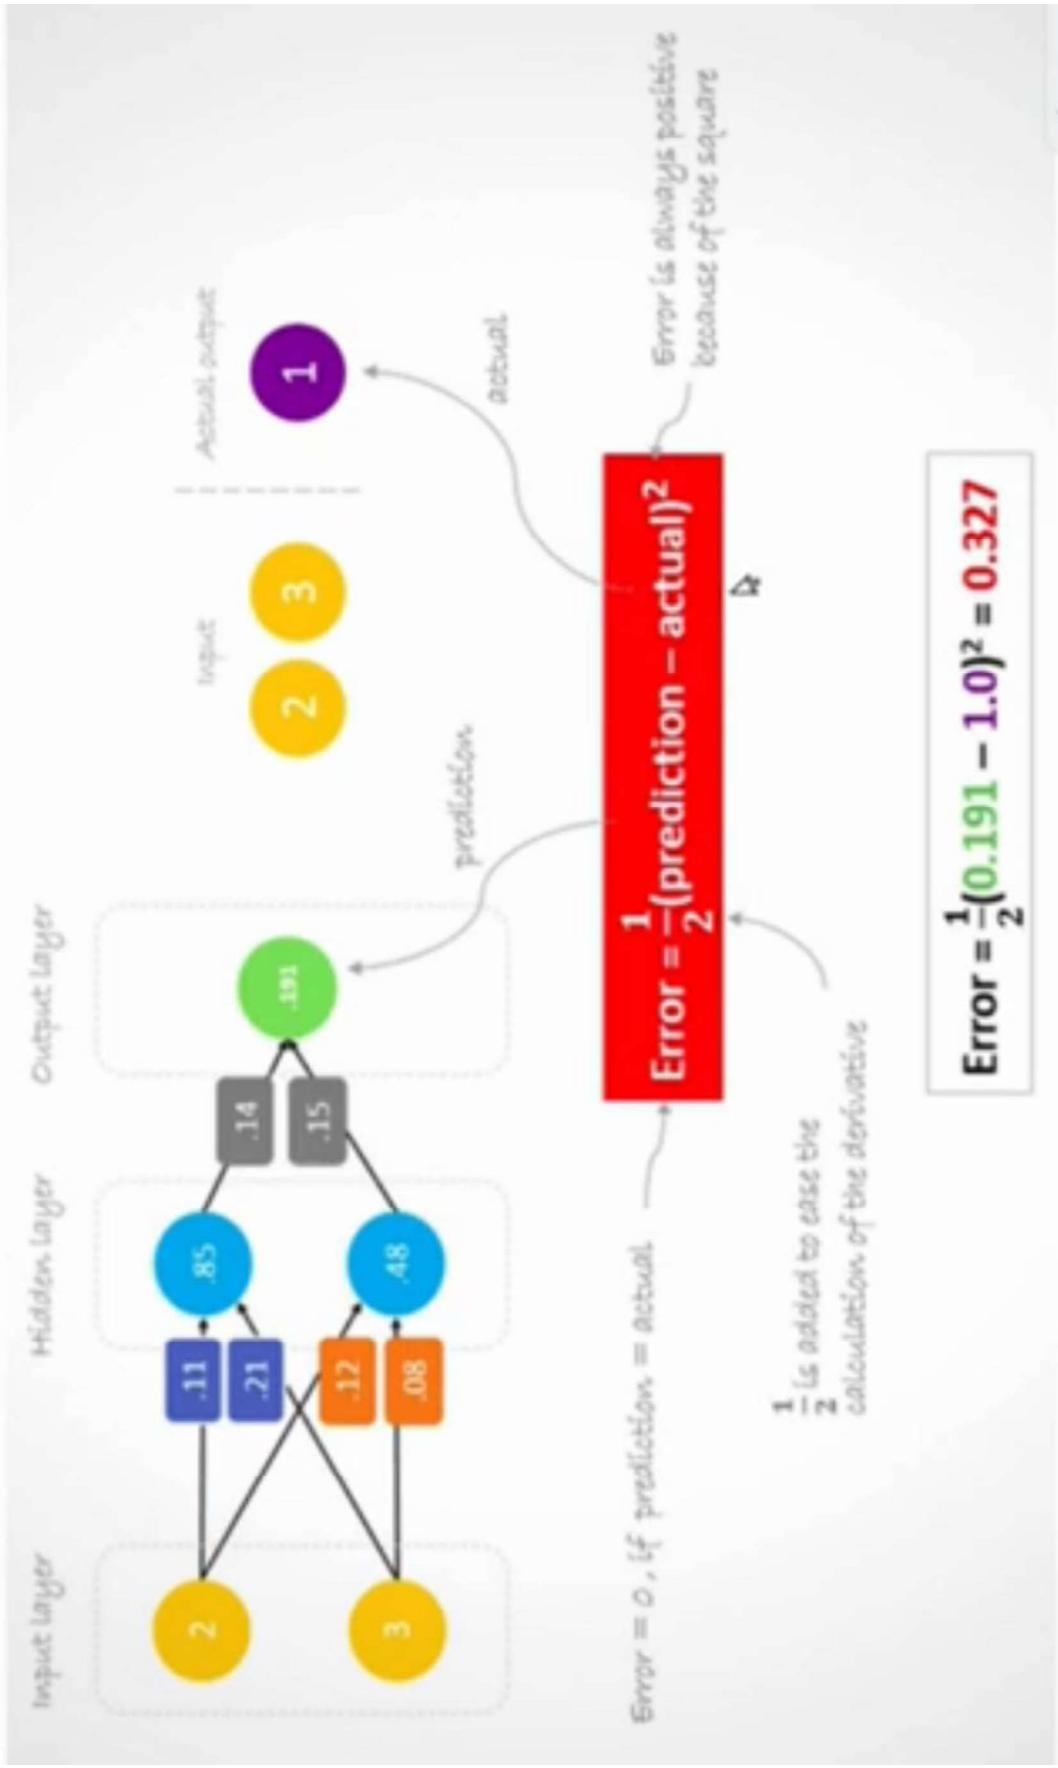
In the backward pass, partial derivatives of the cost function with respects to the different parameters are propagated back through the network.

The process continues until the **error is at the lowest value**.

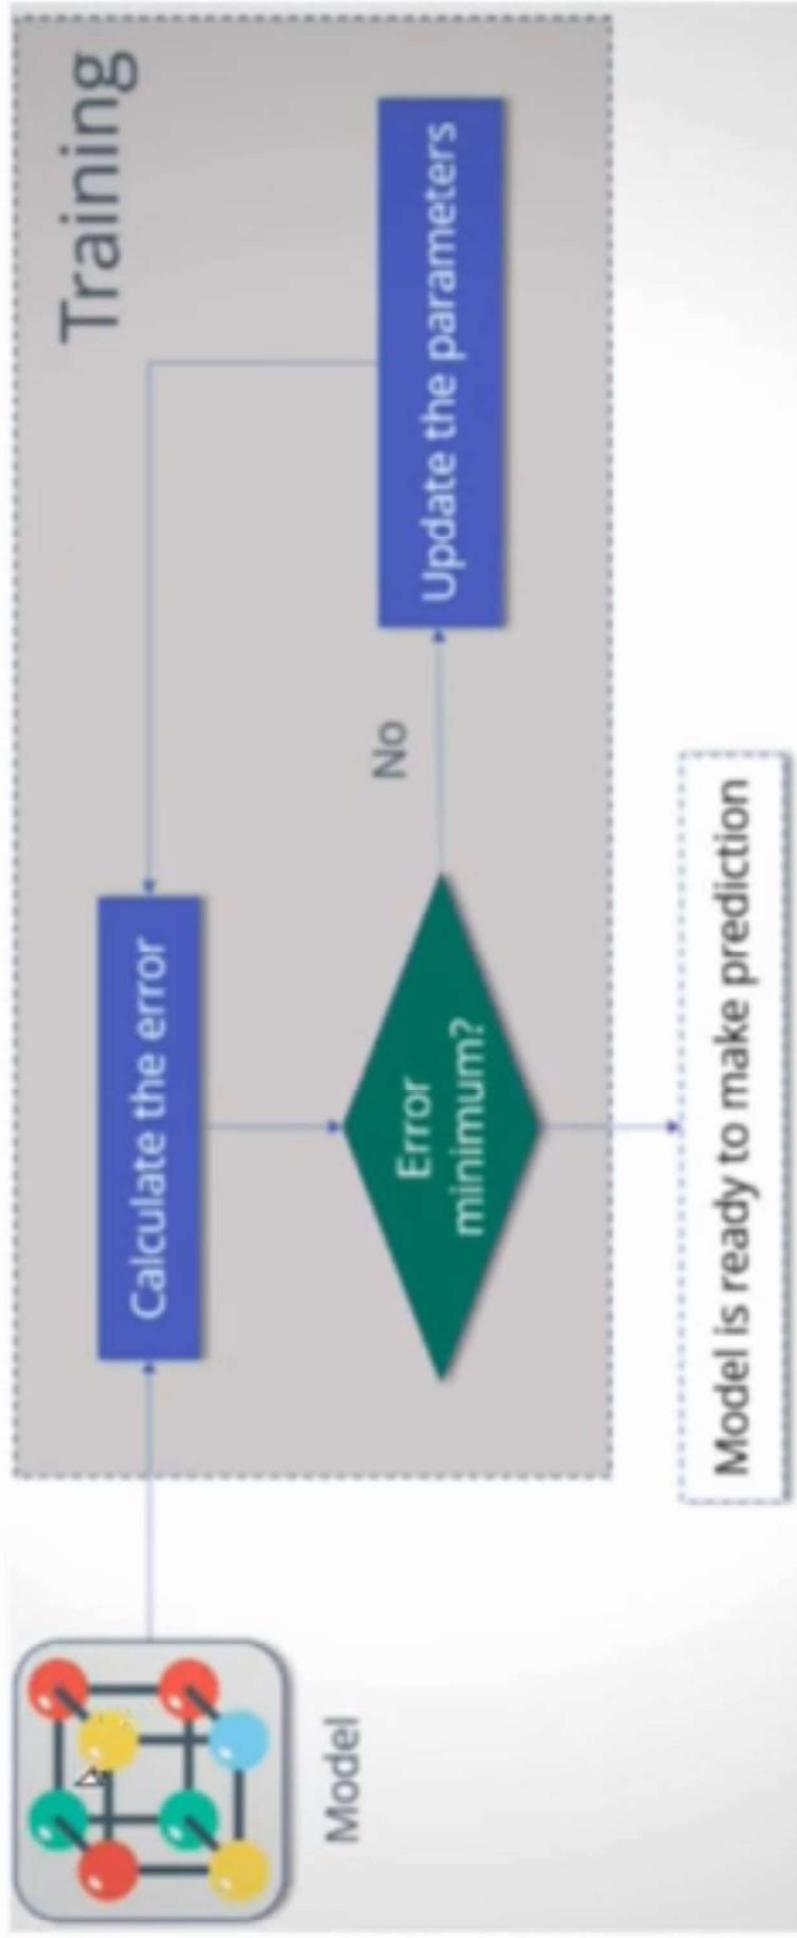
Backpropagation Steps

- Step 1: Determine the architecture
 - how many input and output neurons; what output encoding?
 - hidden neurons and layers.
- Step 2: Initialize all weights and biases to small random values, typically $\in [-1,1]$, choose a learning rate η .
- Step 3: Repeat until termination criteria satisfied
 - Present a training example and propagate it through the network (forward pass)
 - Calculate the actual output
 - Inputs applied
 - Multiplied by weights
 - Summed
 - 'Squashed' by sigmoid activation function
 - Output passed to each neuron in next layer
 - Adapt weights starting from the output layer and working backwards (backward pass)

Why we need Backpropagation?

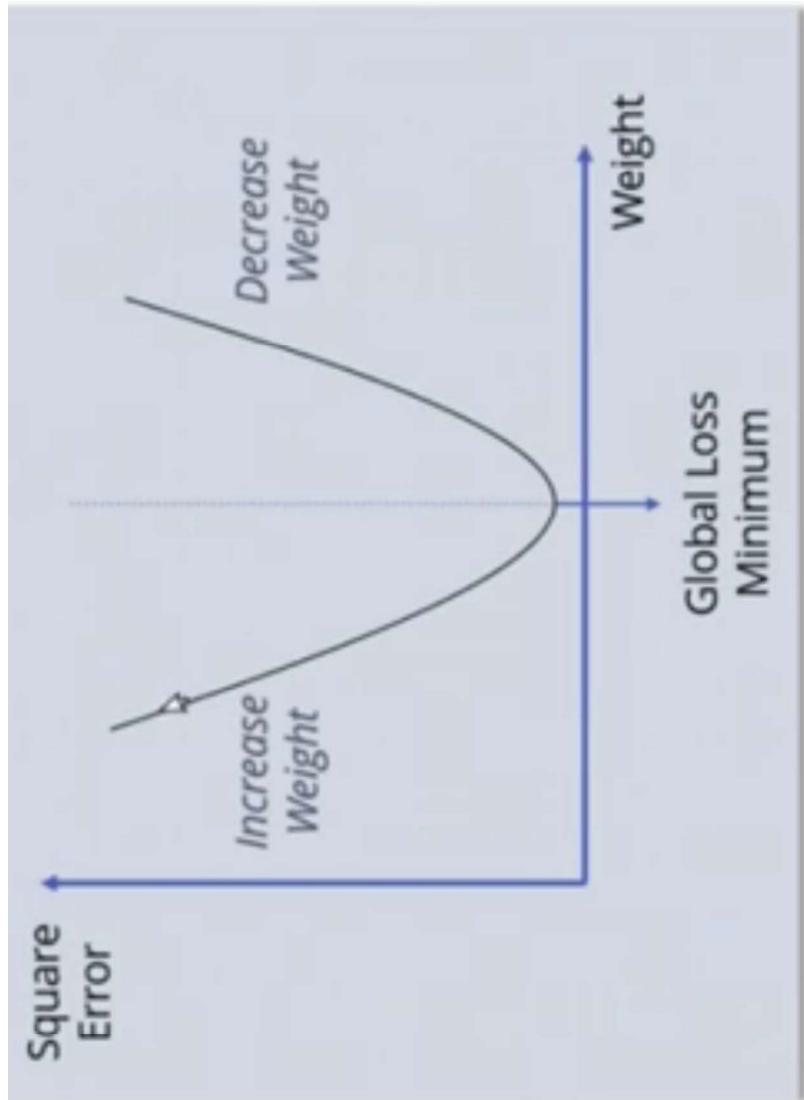


OVERVIEW:



Gradient Descent

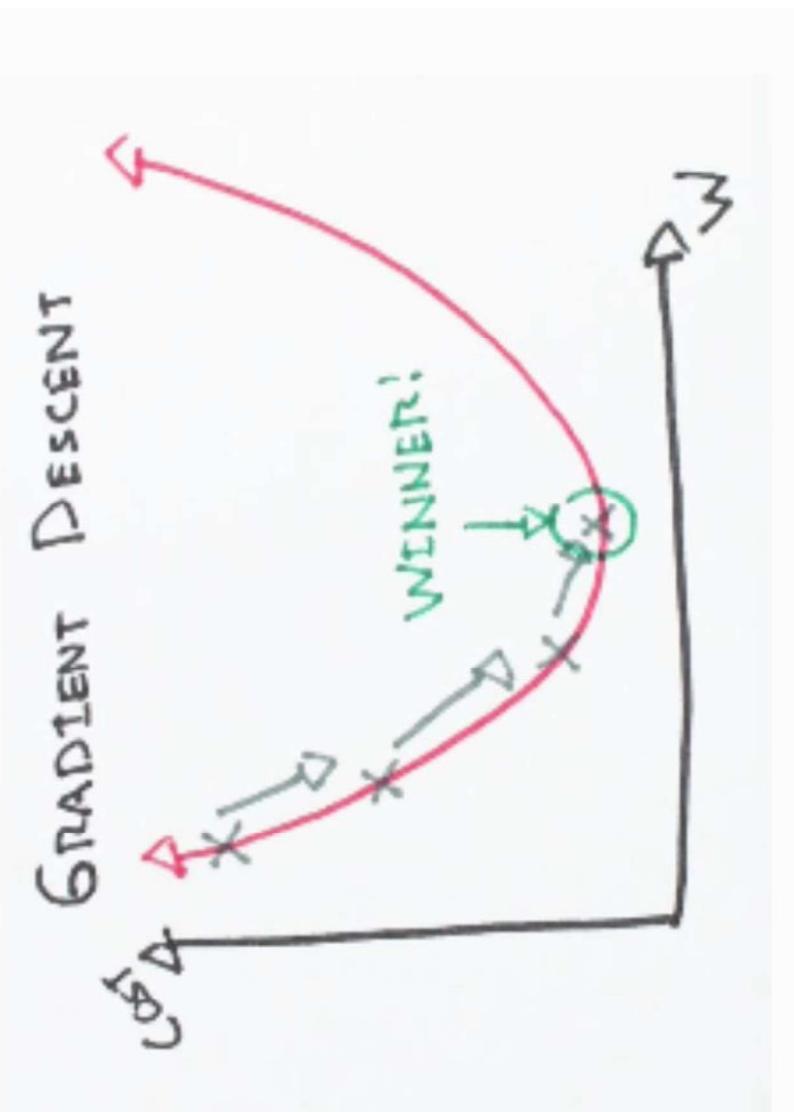
Gradient Descent is known as one of the most commonly used optimization algorithms to train machine learning models by means of minimizing errors between actual and expected results. Further, gradient descent is also used to train Neural Networks.



- Update the weights using gradient descent.
- Gradient descent is used for finding the minimum of a function.
- In our case we want to minimize the error function.

Gradient Descent

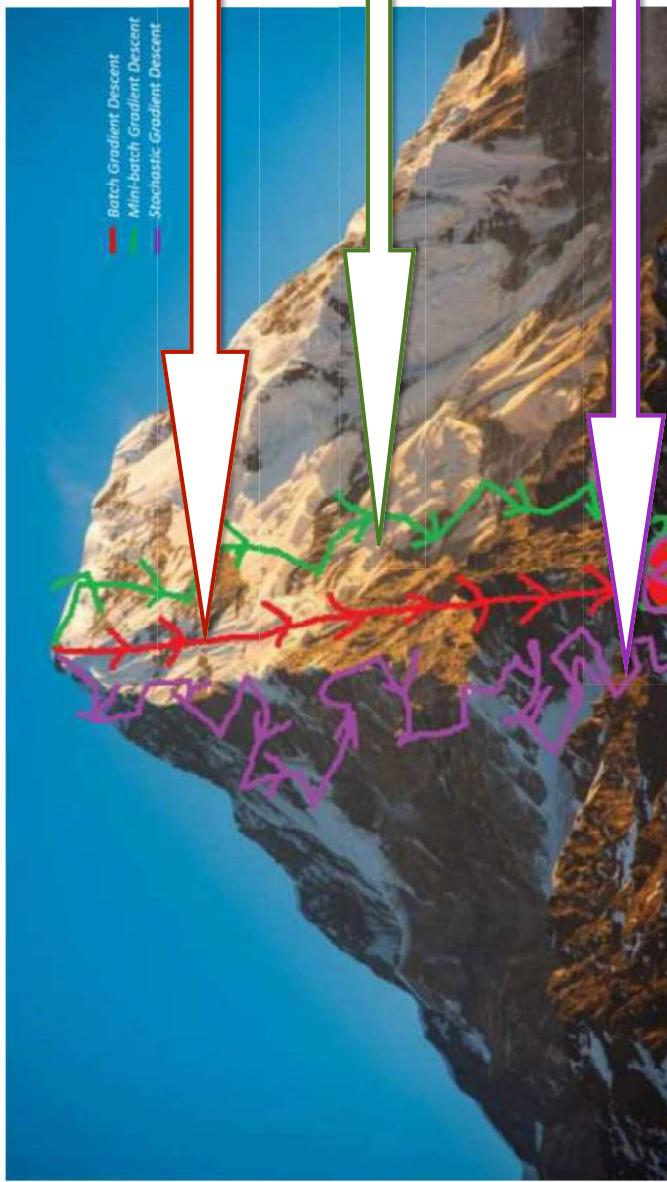
Gradient descent is an optimization algorithm used to minimize some function by iteratively moving in the direction of steepest descent as defined by the negative of the gradient. In machine learning, we use gradient descent to update the **parameters** of our model. Parameters refer to coefficients in **Linear Regression** and **weights** in neural networks.



GD optimiser, not used in Deep Learning (for large dataset)

- GD means considering all the data points from the dataset.[weights get converge quickly].
- **Disadvantage:** Suppose data points is more than million, that point it will need more time to load it and computational power will be more, more number of resources required more. (because of million record).
- To overcome this, SGD used (more specifically mini batch SGD-taken k- data points). Here, try to select batch, batch may be 100/200 data points depends.
- In SGD, it tries to update weight for each and every data points.
- So, takes more time to reach global minima. To remove noise in SGD—momentum is used.

Gradient Descent



n data point

Batch Gradient Descent

K data point

Mini Batch Gradient Descent

1 data point

Stochastic Gradient Descent

SGD: Update the weights for each and every datapoint.

Gradient descent is an iterative optimization algorithm for finding the minimum of a function; in our case we want to minimize the error function.

- Batch gradient descent refers to calculating the derivative from all training data before calculating an update.
- Stochastic gradient descent refers to calculating the derivative from each training data instance and calculating the update immediately.

Stochastic Gradient Descent

$$\omega_{new} = \omega_{old} - \eta \times \frac{\partial L}{\partial w_{old}}$$
$$d_{loss} = \sum_{i=1}^K (y_i - \hat{y}_i)^2 \rightarrow \text{SGD}$$
$$= \sum_{i=1}^K (y_i - \hat{y}_i)^2 \rightarrow \text{SGD}$$
$$= (y - \hat{y})^2 \rightarrow \text{SGD}$$

Derivative of loss w.r.t derivative of w_t

if I am finding derivative Slope for all data points.

↑ n Data Points → Gradient Descent

↓ 1-D.P. → S.G.D. →

K Data Points (where $K < n$) → mini Batch S.G.D.

↓

Diagram illustrating the update step for a single data point:

$$w_t \leftarrow w_t - \eta \cdot \frac{\partial L}{\partial w_t}$$

where $L = \frac{1}{2} (y_i - \hat{y}_i)^2$

Here, Slope for Single data point

for only k points ($k < n$)

for mini batch

Loss = $\sum_{i=1}^{K(n)} (y_i - \hat{y}_i)^2 \rightarrow$ Mini Batch S.G.D.

$\frac{\partial L}{\partial w_t} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i) \rightarrow$ Gr.D.

$= (y_0 - \hat{y}_0)^2 + (y_1 - \hat{y}_1)^2 + \dots + (y_n - \hat{y}_n)^2 \rightarrow S.G.D.$

Differentiation Rules

Constant Rule

$$\frac{d}{dx}[c] = 0$$

Power Rule

$$\frac{d}{dx}x^n = nx^{n-1}$$

Product Rule

$$\frac{d}{dx}[f(x)g(x)] = f'(x)g(x) + f(x)g'(x)$$

Quotient Rule

$$\frac{d}{dx}\left[\frac{f(x)}{g(x)}\right] = \frac{g(x)f'(x) - f(x)g'(x)}{\left[g(x)\right]^2}$$

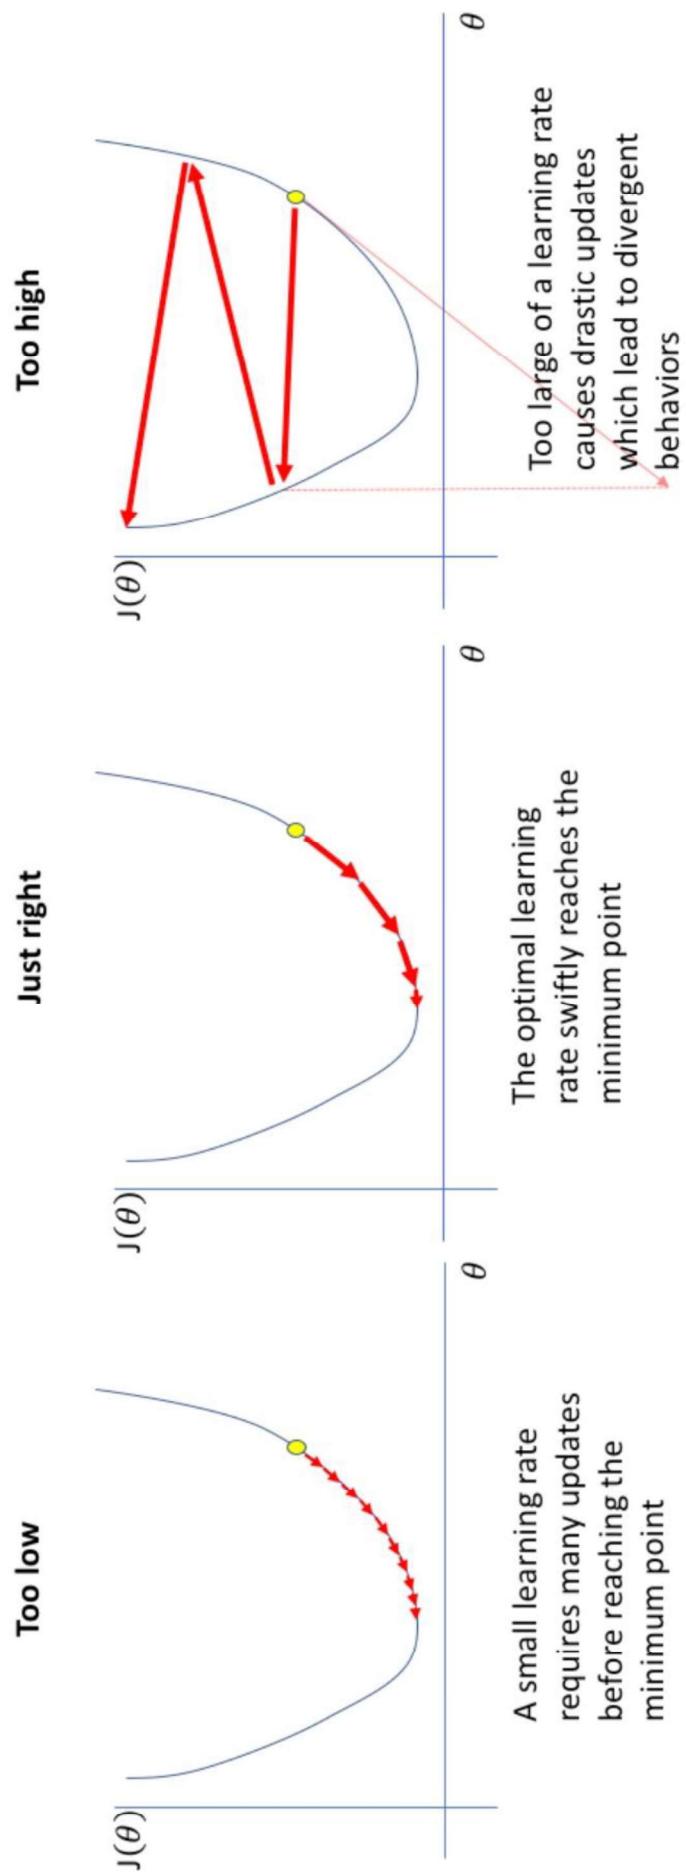
Chain Rule

$$\frac{d}{dx}[f(g(x))] = f'(g(x))g'(x)$$

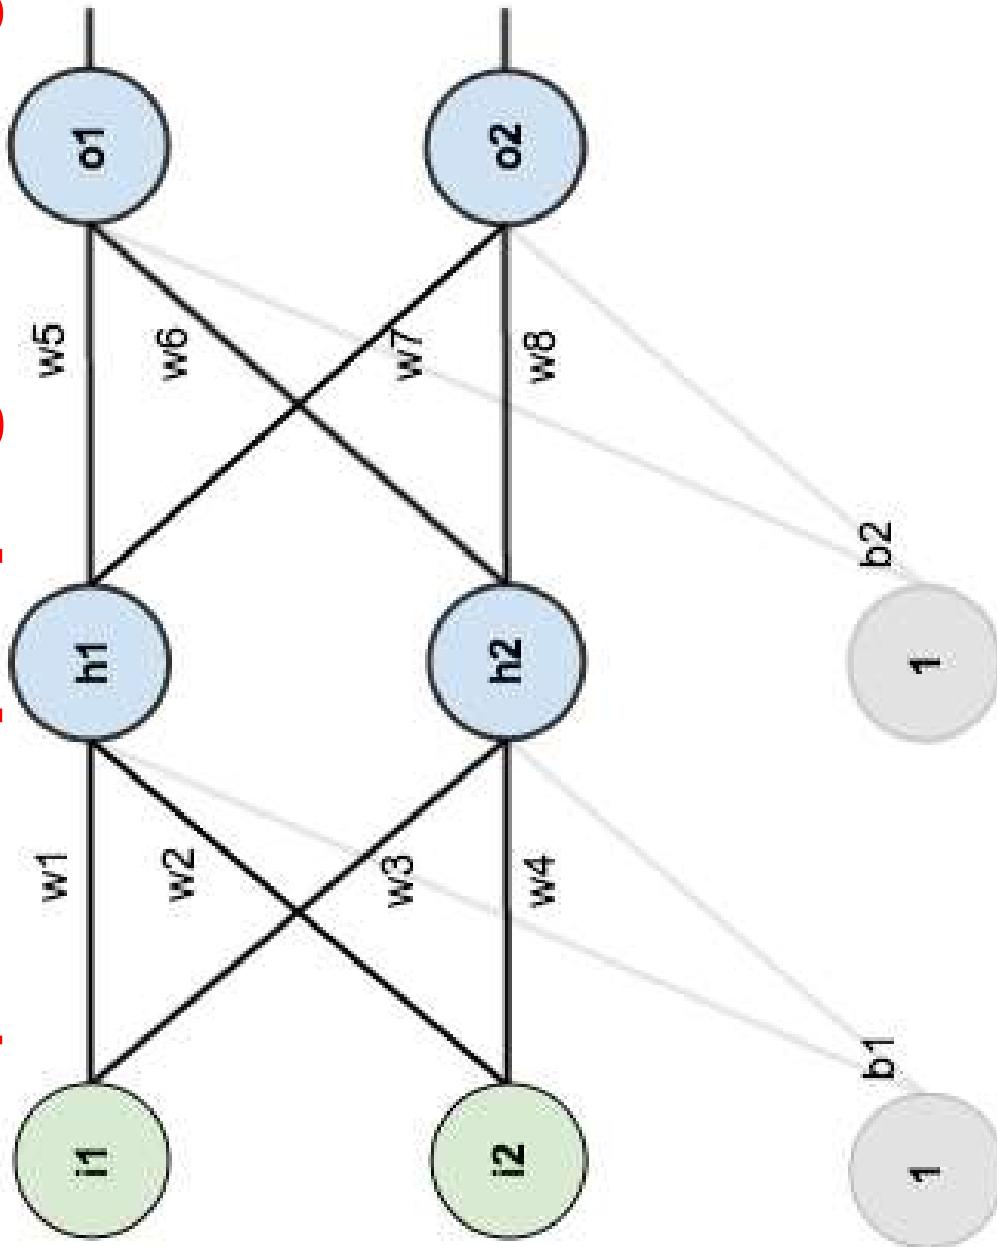
Learning Rate

- The amount that the weights are updated during training is referred to as the step size or the “*learning rate*.”
- Specifically, the learning rate is a configurable hyperparameter used in the training of neural networks that has a small positive value, often in the range between 0.0 and 1.0.
- ... *learning rate*, a positive scalar determining the size of the step.
- The learning rate is often represented using the notation of the lowercase Greek letter eta (η).
- During training, the backpropagation of error estimates the amount of error for which the weights of a node in the network are responsible. Instead of updating the weight with the full amount, it is scaled by the learning rate.
- This means that a learning rate of 0.1, a traditionally common default value, would mean that weights in the network are updated $0.1 * (\text{estimated weight error})$ or 10% of the estimated weight error each time the weights are updated.

When the learning rate is too large, gradient descent can inadvertently increase rather than decrease the training error. [...] When the learning rate is too small, training is not only slower, but may become permanently stuck with a high training error.



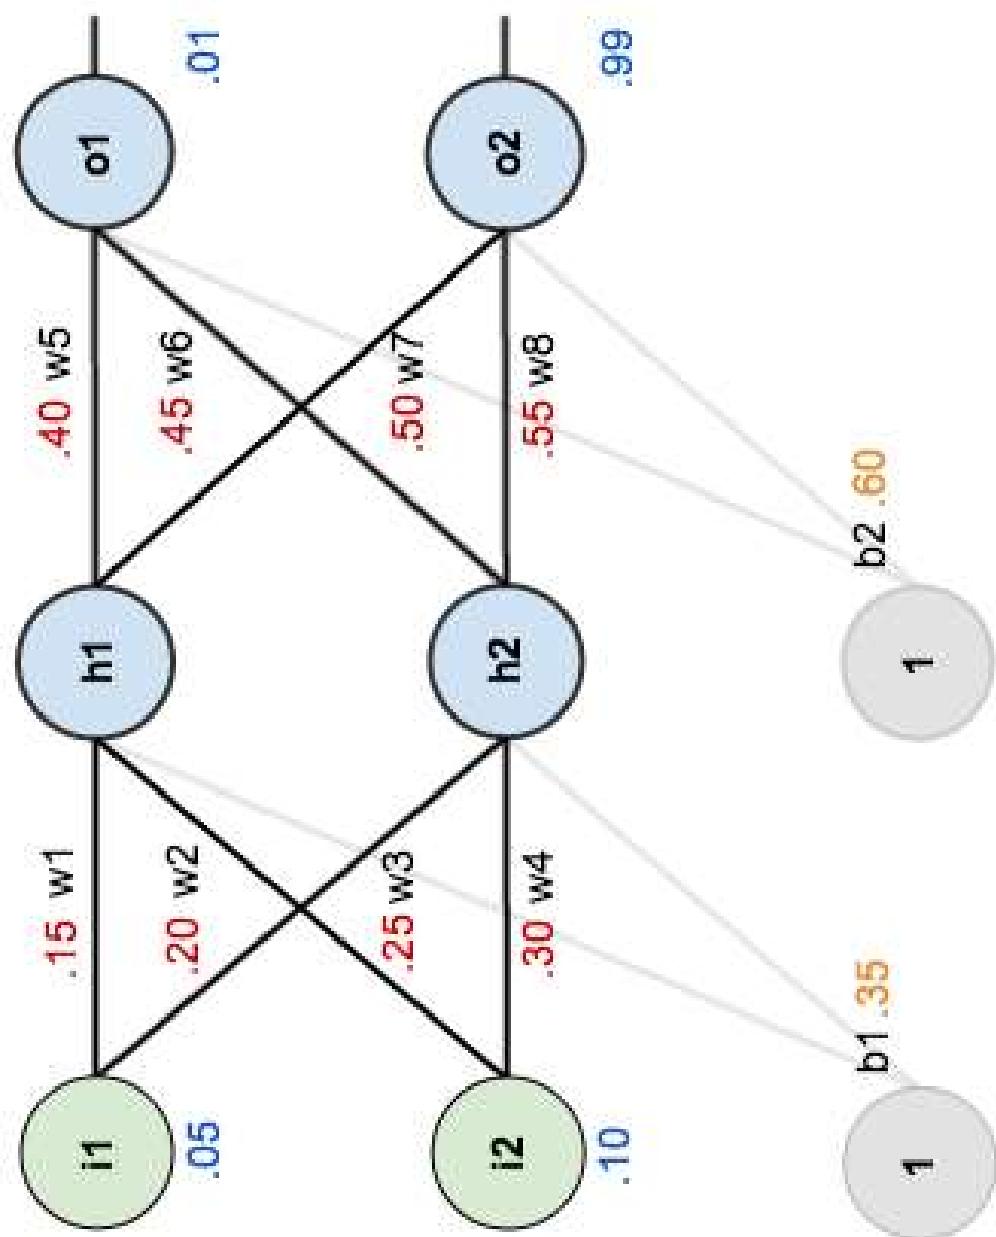
Step-by-Step Backpropagation Algorithm



The goal of backpropagation is to optimize the weights so that the neural network can learn how to correctly map arbitrary inputs to outputs.

For the rest of this tutorial we're going to work with a single training set: given inputs 0.05 and 0.10, we want the neural network to output 0.01 and 0.99.

$i_1=0.05, i_2=0.10$
 $w_1=0.15, w_2=0.20$
 $w_3=0.25, w_4=0.30$
 $b_1=0.3$
 $w_5=0.4, w_6=0.45$
 $w_7=0.5, w_8=0.55$
 $b_2=0.6$
 $o_1=0.01, o_2=0.99$



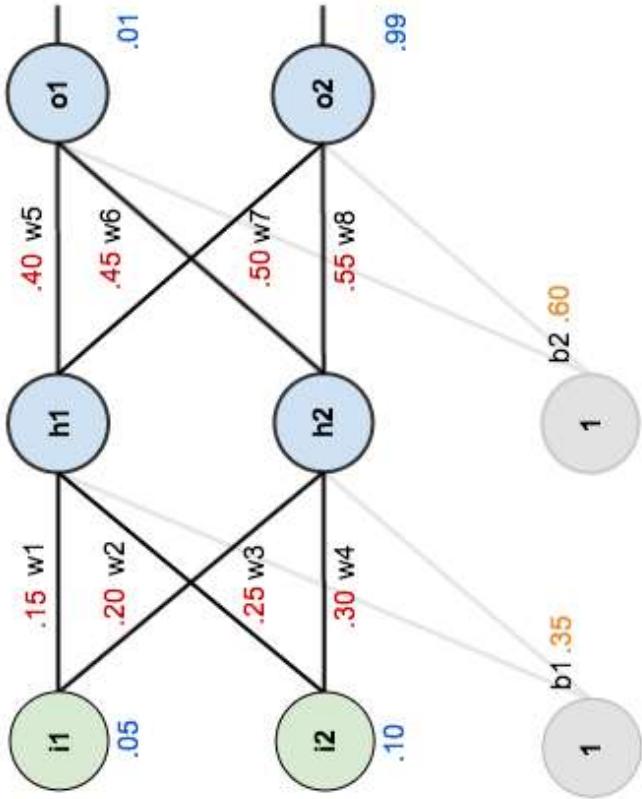
Forward Pass

To begin, let's see what the neural network currently predicts given the weights and biases above and inputs of 0.05 and 0.10. To do this we'll feed those inputs forward through the network.

$$net_{h1} = w_1 * i_1 + w_2 * i_2 + b_1 * 1$$

$$net_{h1} = 0.15 * 0.05 + 0.2 * 0.1 + 0.35 * 1 = 0.3775$$

Squash it using logistic function to get output



$$out_{h1} = \frac{1}{1+e^{-net_{h1}}} = \frac{1}{1+e^{-0.3775}} = 0.593269992$$

Same process for h_2

$$out_{h2} = 0.596884378$$

Forward Pass

Now calculating o_1

$$net_{o1} = w_5 * out_{h1} + w_6 * out_{h2} + b_2 * 1$$

$$net_{o1} = 0.4 * 0.593269992 + 0.45 * 0.596884378 + 0.6 * 1 = 1.105905967$$

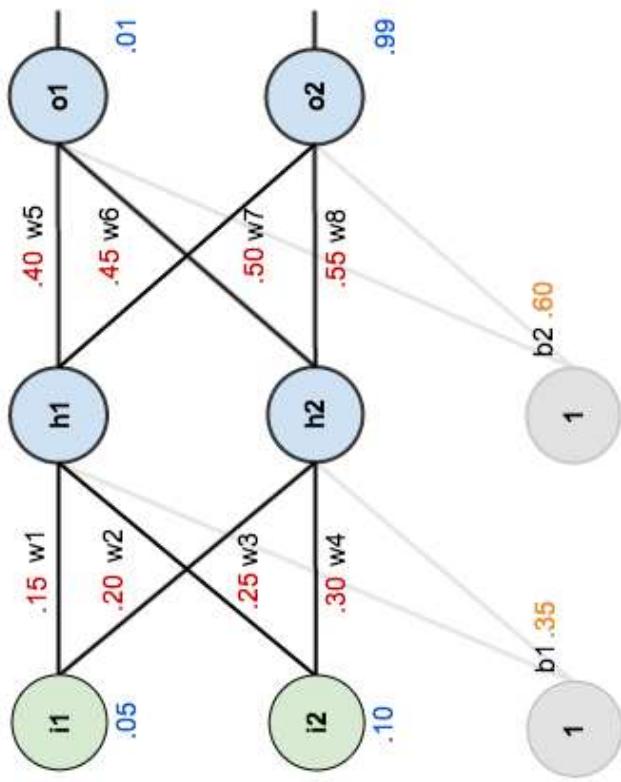
$$out_{o1} = \frac{1}{1+e^{-net_{o1}}} = \frac{1}{1+e^{-1.105905967}} = 0.75136507$$

And carrying out the same process for o_2 we get:

$$out_{o2} = 0.772928465$$

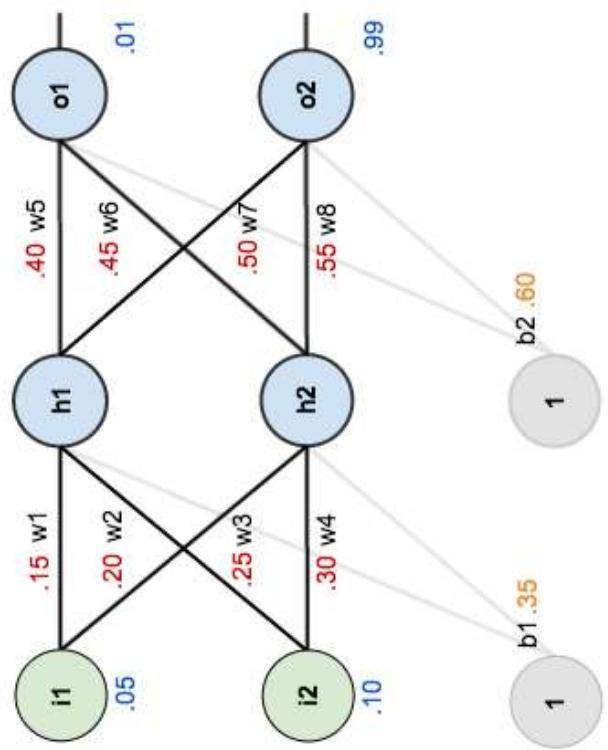
But, Target Values

$$\begin{array}{l} T_1 \\ T_2 \end{array}$$



Calculating Total Error

$$E_{total} = \sum \frac{1}{2}(target - output)^2$$

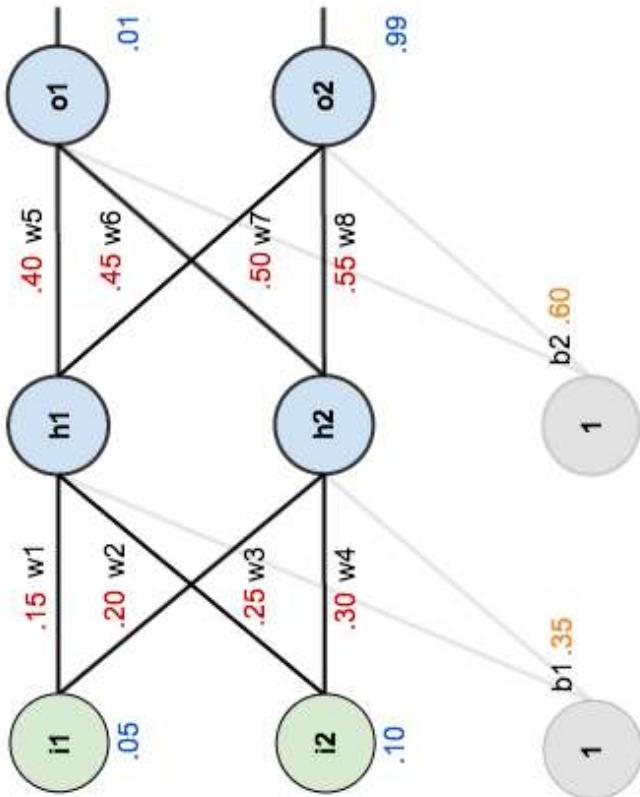


$$E_{o1} = \frac{1}{2}(target_{o1} - output_{o1})^2 = \frac{1}{2}(0.01 - 0.75136507)^2 = 0.274811083$$

$$E_{o2} = 0.023560026$$

$$E_{total} = E_{o1} + E_{o2} = 0.274811083 + 0.023560026 = 0.298371109$$

Backward Pass



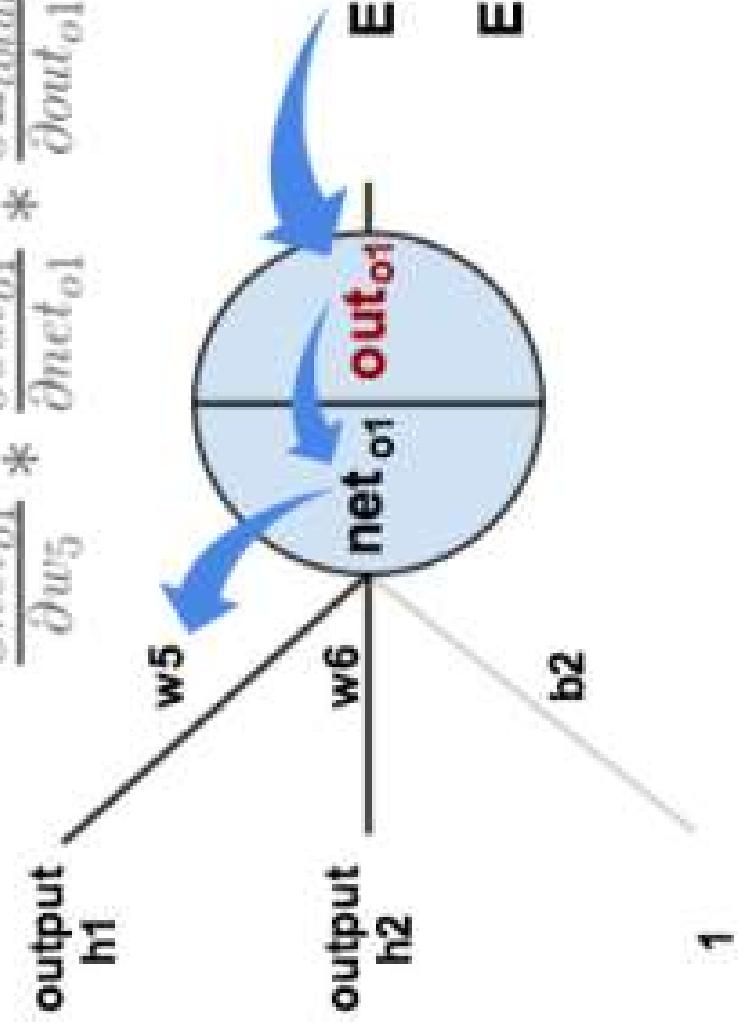
Our goal with backpropagation:

- Is to update each of the weights in the network using gradient descent.
- So that they cause the actual output to be closer to the target output.
- Thereby minimizing the error for each output neuron and the network as a whole.

Chain Rule

To calculate E_{out} at w_5

$$\frac{\partial net_{o1}}{\partial w_5} * \frac{\partial out_{o1}}{\partial net_{o1}} * \frac{\partial E_{total}}{\partial out_{o1}} = \frac{\partial E_{total}}{\partial w_5}$$



$$E_{o1} = \frac{1}{2}(\text{target}_{o1} - \text{out}_{o1})^2$$

$$E_{\text{total}} = E_{o1} + E_{o2}$$

Output Layer.

To Update Weights:

OUTPUT LAYER:

Consider w5.

$$E_{total} = E_{o1} + E_{o2}$$

$$E_{o1} = \frac{1}{2}(target_{o1} - out_{o1})^2$$

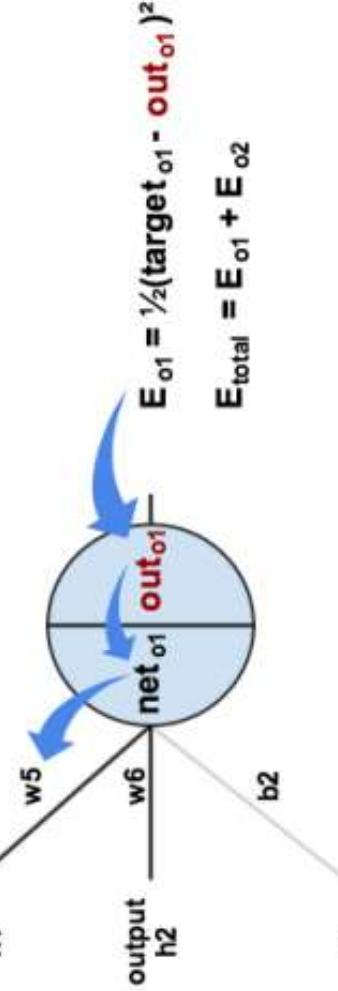
$$out_{o1} = \frac{1}{1+e^{-net_{o1}}}$$

$$net_{o1} = w_5 * out_{h1} + w_6 * out_{h2} + b_2 * 1$$

$$\frac{\partial E_{total}}{\partial w_5} \quad \text{partial derivative of } E_{total} \text{ wrt w5.}$$

$$\frac{\partial E_{total}}{\partial w_5} = \frac{\partial E_{total}}{\partial out_{o1}} * \frac{\partial out_{o1}}{\partial net_{o1}} * \frac{\partial net_{o1}}{\partial w_5}$$

$$\frac{\partial net_{o1}}{\partial w_5} * \frac{\partial out_{o1}}{\partial net_{o1}} * \frac{\partial E_{total}}{\partial out_{o1}} = \frac{\partial E_{total}}{\partial w_5}$$



$$E_{o1} = \frac{1}{2}(target_{o1} - out_{o1})^2$$

$$E_{total} = E_{o1} + E_{o2}$$

$$\frac{\partial E_{total}}{\partial w_5} = \frac{\partial E_{total}}{\partial out_{o1}} * \frac{\partial out_{o1}}{\partial net_{o1}} * \frac{\partial net_{o1}}{\partial w_5}$$

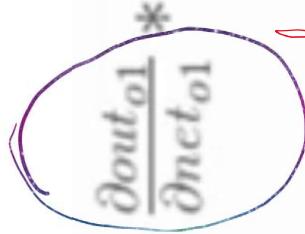
Apply P. & Q
on both side w.r.t

$$E_{total} = \frac{1}{2}(target_{o1} - out_{o1})^2 + \frac{1}{2}(target_{o2} - out_{o2})^2$$

$$\frac{\partial E_{total}}{\partial out_{o1}} = 2 * \frac{1}{2}(target_{o1} - out_{o1})^{2-1} * -1 + 0$$

$$\frac{\partial E_{total}}{\partial out_{o1}} = -(target_{o1} - out_{o1}) = -(0.01 - 0.75136507) = 0.74136507$$

$$\frac{\partial E_{total}}{\partial w_5} = \frac{\partial E_{total}}{\partial out_{o1}} * \frac{\partial out_{o1}}{\partial net_{o1}} * \frac{\partial net_{o1}}{\partial w_5}$$



$$out_{o1} = \frac{1}{1 + e^{-net_{o1}}} \quad \text{P.d.}$$

$$out_{o1} = out_{o1}(1 - out_{o1})$$

$$\frac{\partial out_{o1}}{\partial net_{o1}} = f'(x) = \frac{e^x}{(1 + e^x)^2} = \frac{e^x}{(1 + e^x)^2} = f(x)(1 - f(x)).$$

$$\frac{\partial out_{o1}}{\partial net_{o1}} = 0.75136507(1 - 0.75136507) = 0.186815602$$

$$\frac{\partial E_{total}}{\partial w_5} = \frac{\partial E_{total}}{\partial out_{o1}} * \frac{\partial out_{o1}}{\partial net_{o1}} * \frac{\partial net_{o1}}{\partial w_5}$$

$$net_{o1} = w_5 * out_{h1} + \underbrace{w_6 * out_{h2} + b_2 * 1}_{\text{Change in } w_5}$$

$$\frac{\partial net_{o1}}{\partial w_5} = 1 * out_{h1} * w_5^{(1-1)} + 0 + 0 = out_{h1} = 0.593269992$$

Putting it all together

$$\frac{\partial E_{total}}{\partial w_5} = \frac{\partial E_{total}}{\partial out_{o1}} * \frac{\partial out_{o1}}{\partial net_{o1}} * \frac{\partial net_{o1}}{\partial w_5}$$

$$\frac{\partial E_{total}}{\partial w_5} = 0.74136507 * 0.186815602 * 0.59326992 = 0.082167041$$

Change in w_5

- Update the weight.
- To decrease the error, we then subtract this value from the current weight.

$$w_5^+ = w_5 - \eta * \frac{\partial E_{total}}{\partial w_5} = 0.4 - 0.5 * 0.082167041 = 0.35891648$$

Some way

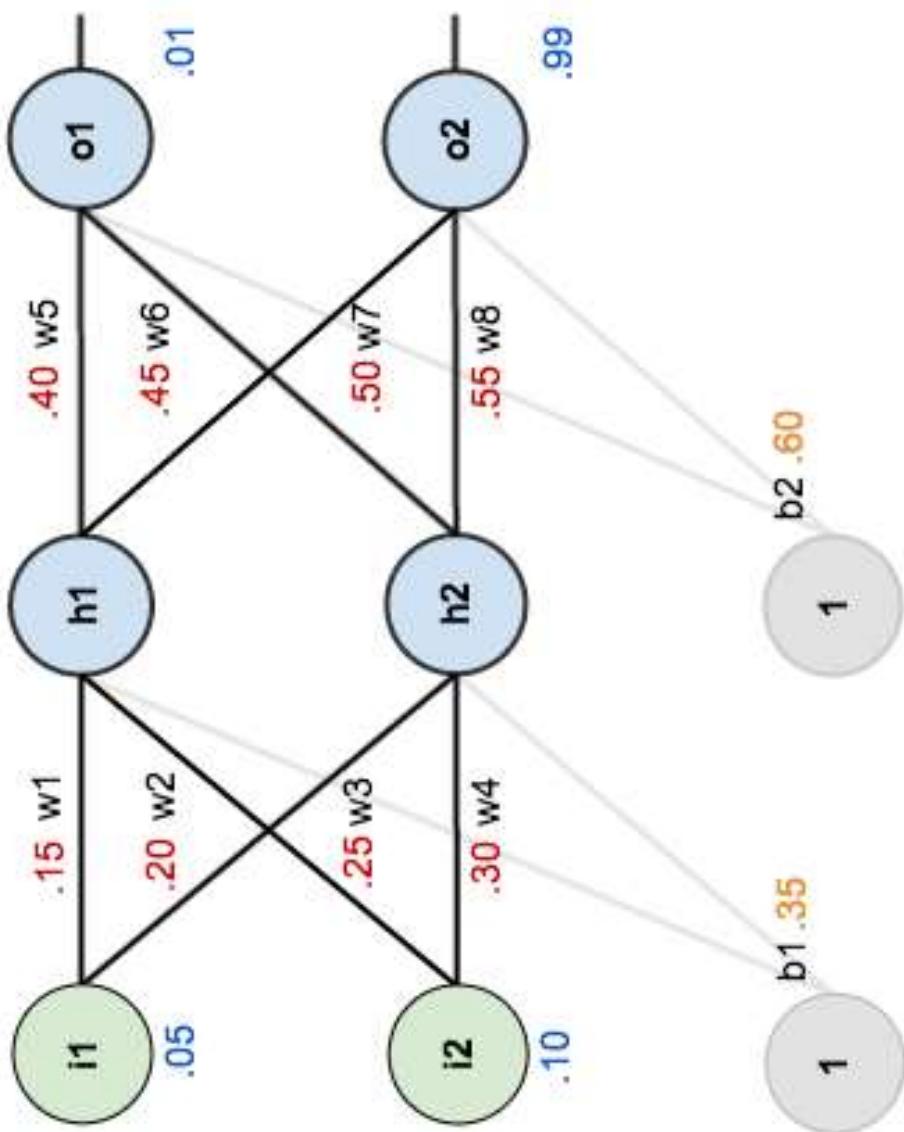
η = learning rate (eta)

$$w_6^+ = 0.408666186$$

$$w_7^+ = 0.511301270$$

$$w_8^+ = 0.561370121$$

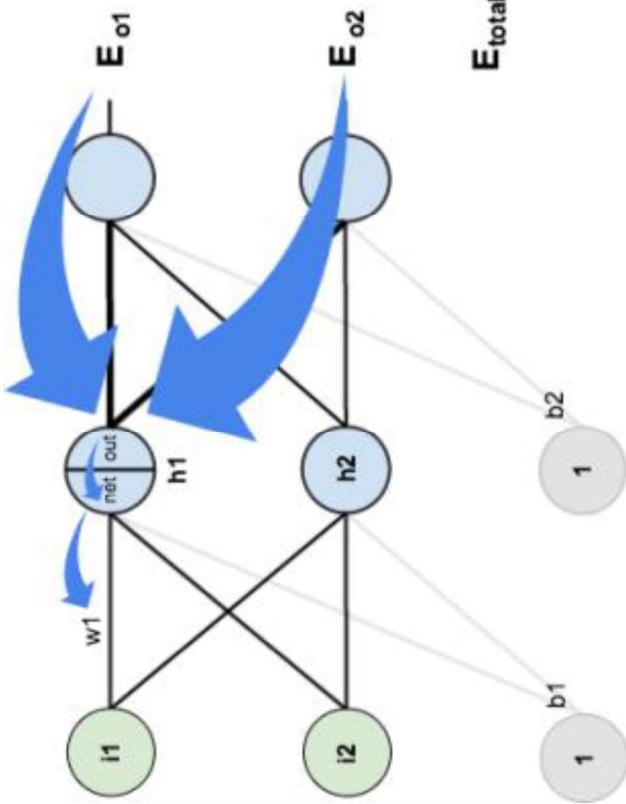
We perform the *actual* updates in the neural network after we have the new weights leading into the hidden layer neurons (ie, we use the original weights, not the updated weights, when we continue the backpropagation algorithm below).



Hidden Layer:

Next, we'll continue the backwards pass by calculating new values for

Upgrading w₁
HIDDEN LAYER:
 Consider w₁.



$$\frac{\partial E_{total}}{\partial w_1} = \frac{\partial E_{total}}{\partial out_{h1}} * \frac{\partial out_{h1}}{\partial net_{h1}} * \frac{\partial net_{h1}}{\partial w_1}$$

$$\frac{\partial E_{total}}{\partial out_{h1}} = \frac{\partial E_{o1}}{\partial out_{h1}} + \frac{\partial E_{o2}}{\partial out_{h1}}$$

↓
↓ K.T.

$$E_{o1} = \frac{1}{2}(target_{o1} - out_{o1})^2$$

$$out_{o1} = \frac{1}{1+e^{-net_{o1}}}$$

$$net_{o1} = w_5 * out_{h1} + w_6 * out_{h2} + b_2 * 1$$

$$\frac{\partial E_{o1}}{\partial out_{h1}} = \frac{\partial E_{o1}}{\partial out_{o1}} * \frac{\partial out_{o1}}{\partial net_{o1}} * \frac{\partial net_{o1}}{\partial out_{h1}}$$

$$\frac{\partial net_{o1}}{\partial out_{h1}} = w_5$$

$$\frac{\partial E_{o1}}{\partial out_{h1}} = \frac{\partial E_{o1}}{\partial net_{o1}} * \frac{\partial net_{o1}}{\partial out_{h1}} = 0.138498562 * 0.40 = 0.055399425$$

$$\frac{\partial E_{o2}}{\partial out_{h1}} = -0.019049119$$

$$\frac{\partial E_{total}}{\partial out_{h1}} = \frac{\partial E_{o1}}{\partial out_{h1}} + \frac{\partial E_{o2}}{\partial out_{h1}} = 0.055399425 + -0.019049119 = 0.036350306$$

$$\frac{\partial E_{total}}{\partial w_1} = \frac{\partial E_{total}}{\partial out_{h1}} * \frac{\partial out_{h1}}{\partial net_{h1}} * \frac{\partial net_{h1}}{\partial w_1}$$

$$out_{h1} = \frac{1}{1 + e^{-net_{h1}}}$$

$$\frac{\partial out_{h1}}{\partial net_{h1}} = out_{h1}(1 - out_{h1}) = 0.59326999(1 - 0.59326999) = 0.241300709$$

$$net_{h1} = w_1 * i_1 + w_3 * i_2 + b_1 * 1$$

Putting it all together

$$\frac{\partial E_{total}}{\partial w_1} = \frac{\partial E_{total}}{\partial out_{h1}} * \frac{\partial out_{h1}}{\partial net_{h1}} * \frac{\partial net_{h1}}{\partial w_1} = i_1 = 0.05$$

$$\frac{\partial E_{total}}{\partial w_1} = 0.036350306 * 0.241300709 * 0.05 = 0.000438568$$

We can now update w_1 :

$$w_1^+ = w_1 - \eta * \frac{\delta E_{total}}{\partial w_1} = 0.15 - 0.5 * 0.000438568 = 0.149780716$$

Repeating this for w_2, w_3 , and w_4

$$w_2^+ = 0.19956143$$

$$w_3^+ = 0.24975114$$

$$w_4^+ = 0.29950229$$

- Finally, we've updated all of our weights!
 - When we fed forward the 0.05 and 0.1 inputs originally, the error on the network was 0.298371109.
 - After this first round of backpropagation, the total error is now down to 0.291027924.
 - It might not seem like much, but after repeating this process 10,000 times, for example, the error plummets to 0.0000351085.
- Finally, we've updated all of our weights! When we fed forward the 0.05 and 0.1 inputs originally, the error on the network was 0.298371109. After this first round of backpropagation, the total error is now down to 0.291027924. It might not seem like much, but after repeating this process 10,000 times, for example, the error plummets to 0.0000351085. At this point, when we feed forward 0.05 and 0.1, the two outputs neurons generate 0.015912196 (vs 0.01 target) and 0.984065734 (vs 0.99 target).

Backpropagation Algorithm

BACKPROPAGATION(*training_example*, η , n_{in} , n_{out} , n_{hidden})

- Each training example is a pair of the form (x, t) , where (x) is the vector of network input values, and (t) is the vector of target network output values.
- η is the learning rate (e.g., 0.05).
- n_i , is the number of network inputs,
- n_{hidden} the number of units in the hidden layer, and
- n_{out} the number of output units.
- The input from unit i into unit j is denoted x_{ji} , and the weight from unit i to unit j is denoted w_{ji}

Backpropagation Algorithm

- Create a feed-forward network with n_i inputs, n_{hidden} hidden units, and n_{out} output units.

- Initialize all network weights to small random numbers

- Until the termination condition is met, Do

- For each (x, t) , in training examples, Do

- Propagate the input forward through the network:

1. Input the instance x , to the network and compute the output o_u of every unit u in the network.

- Propagate the errors backward through the network

2. For each network unit k , calculate its error term δ_k

$$\delta_k \leftarrow o_k(1 - o_k)(t_k - o_k)$$

Where

$$\Delta w_{ji} = \eta \delta_j x_{ji}$$

4. Update each network weight w_{ji}

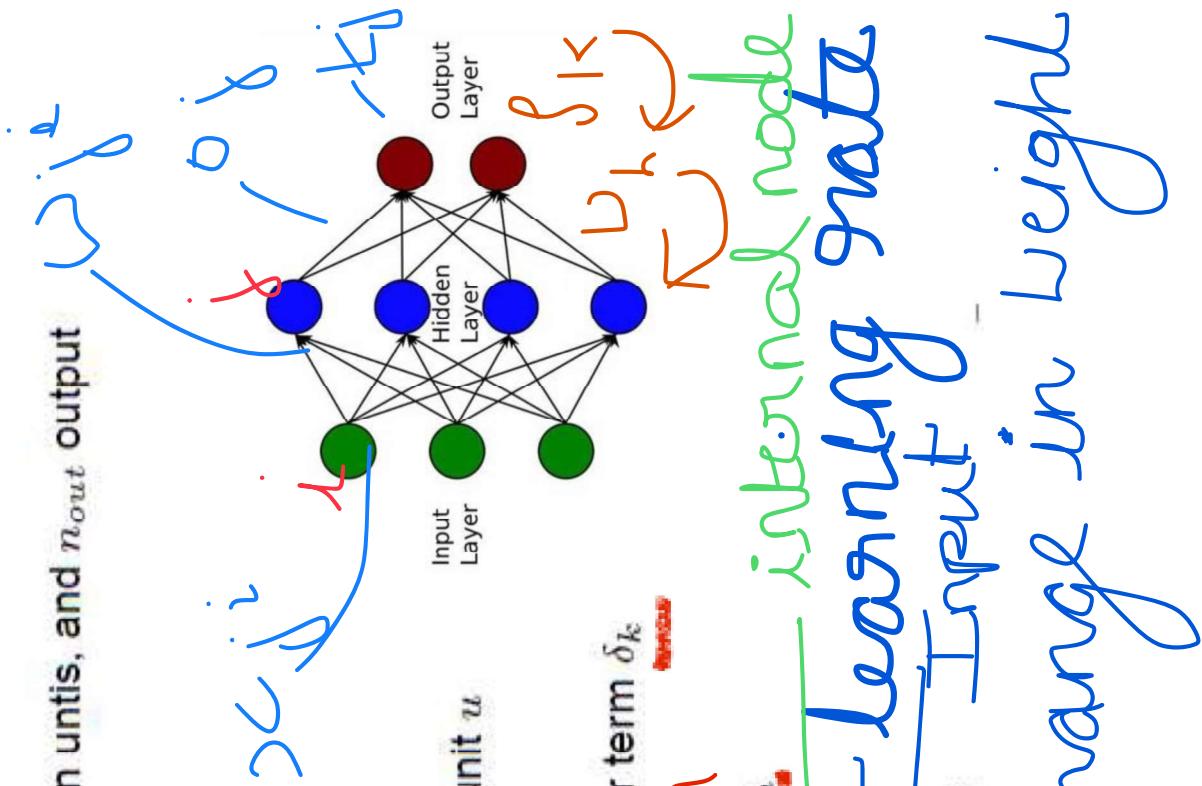
$$w_{ji} \leftarrow w_{ji} + \Delta w_{ji}$$

3. For each network unit h , calculate its error term δ_h

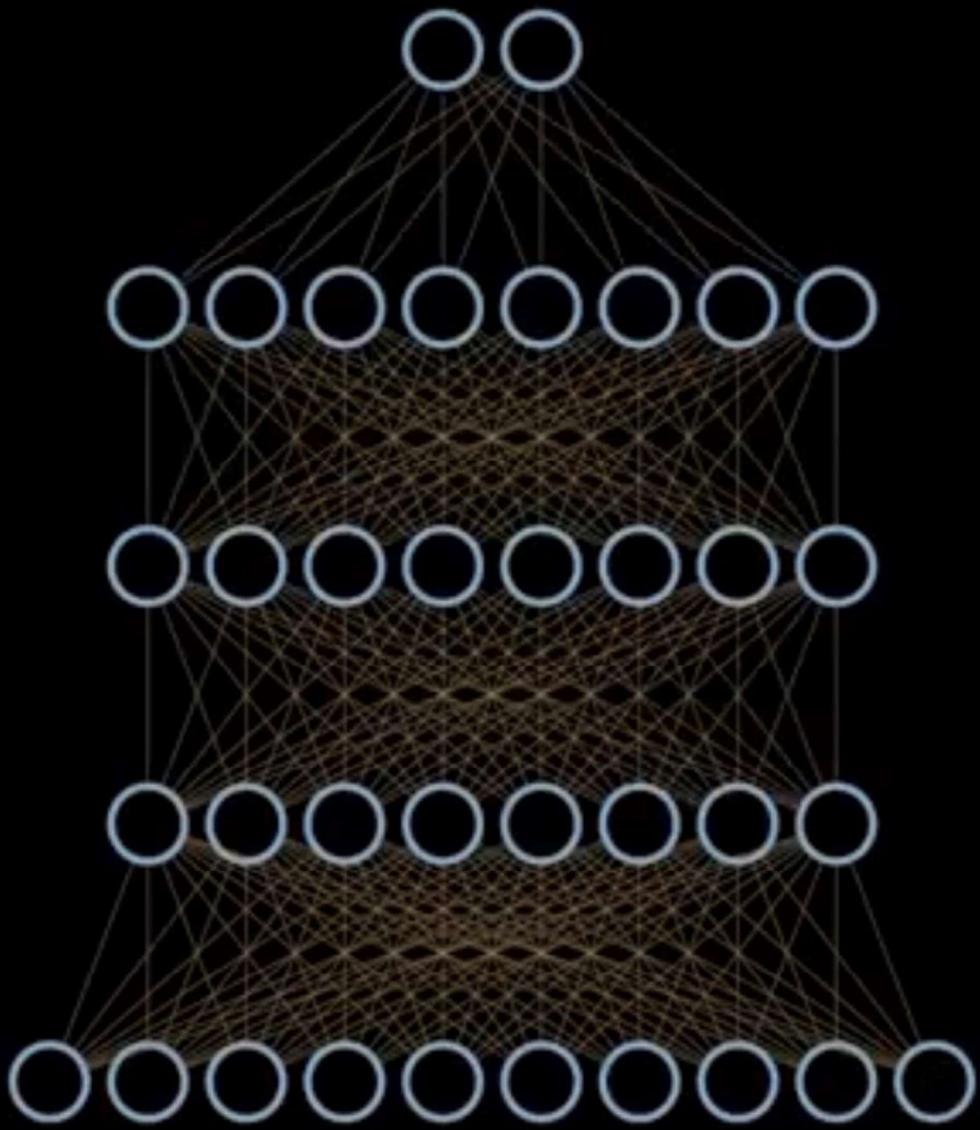
$$\delta_h \leftarrow o_h(1 - o_h) \sum_{k \in \text{outputs}} w_{h,k} \delta_k$$

BACKPROPAGATION Algorithm

- create a feed-forward network with n_{in} inputs, n_{hidden} hidden units, and n_{out} output units
 - Initialize all network weights to small random numbers
 - Until the termination condition is met, Do
 - For each $\langle \vec{x}, \vec{t} \rangle$ in training_examples, Do
 - Propagate the input forward through the network:
1. Input \vec{x} to the network and compute o_u of every unit u
- Propagate the errors back through the network:
2. For each network output unit k , calculate its error term δ_k
$$\delta_k \leftarrow o_k(1 - o_k)(t_k - o_k)$$
 3. For each hidden unit h , calculate its error term δ_h
$$\delta_h \leftarrow o_h(1 - o_h) \sum_{k \in \text{outputs}} w_{kh}\delta_k$$
 4. Update each weight w_{ji} where $\Delta w_{ji} = \eta \delta_j x_{ji}$
- By rule of thumb*
- From δ_k & δ_h*
- Change in weight*



Layer sizes: 10, 8, 8, 8, 2

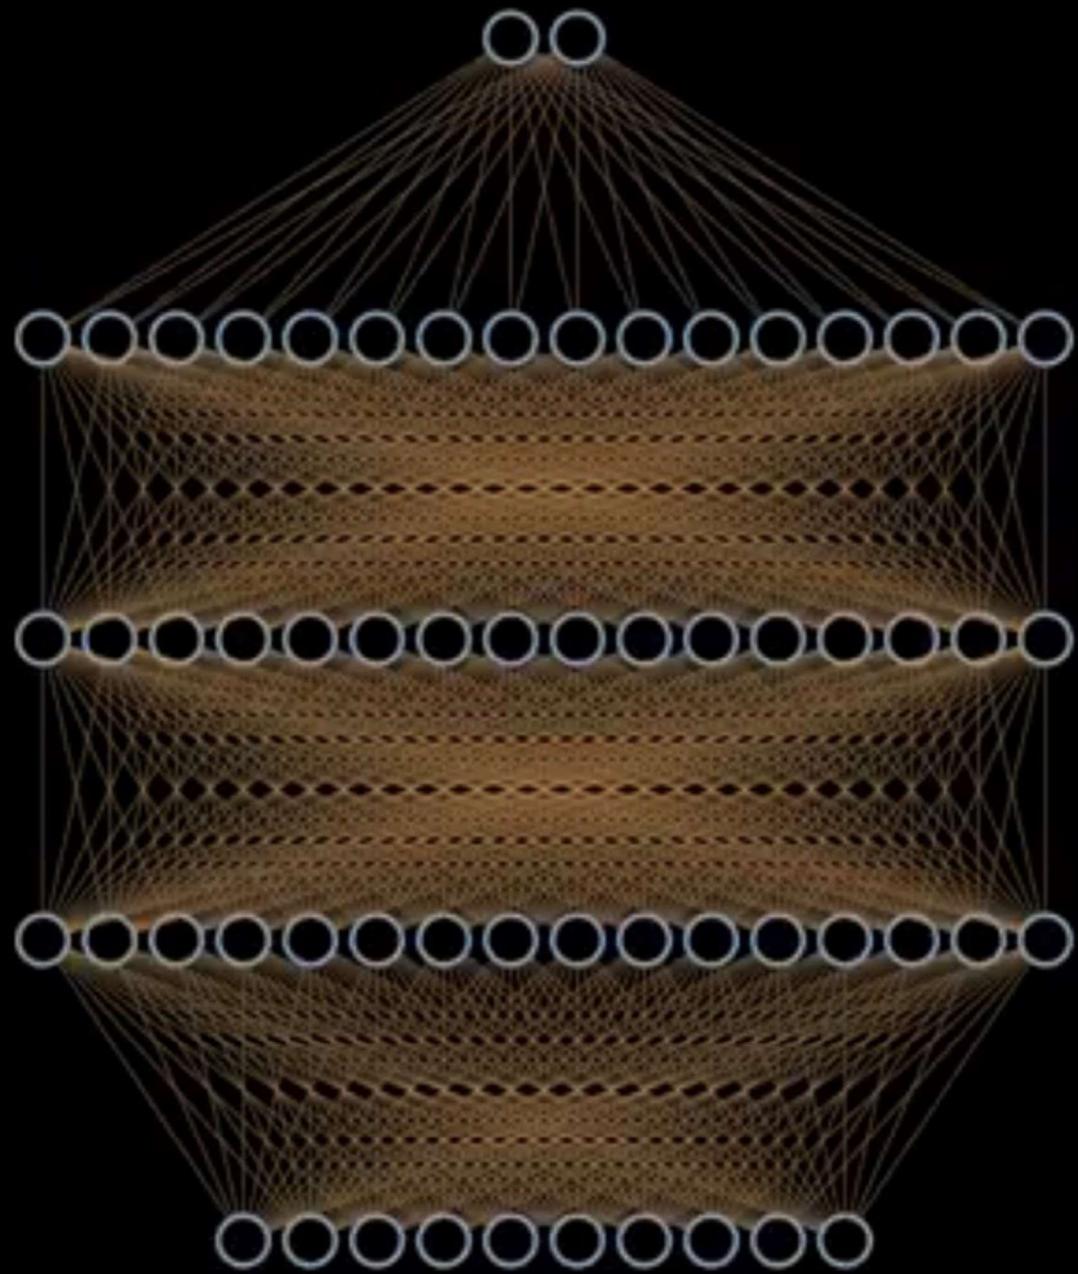


$$\begin{array}{rcl} \text{Weights:} & 224 & \\ \text{Biases:} & 36 & \\ \hline \text{Params:} & 260 & \end{array}$$

<https://nnfs.io>

every neuron is connected to the subsequent layer of neurons in folds.

Layer sizes: 10, 16, 16, 16, 2

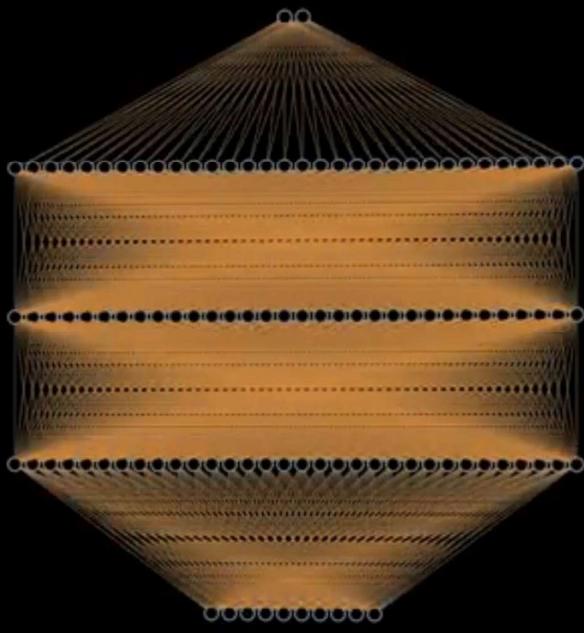


$$\begin{array}{rcl} \text{Weights:} & 704 \\ \text{Biases:} & 60 \\ \hline \text{Params:} & 764 \end{array}$$

<https://nnfs.io>

and then every neuron is a unique bias.

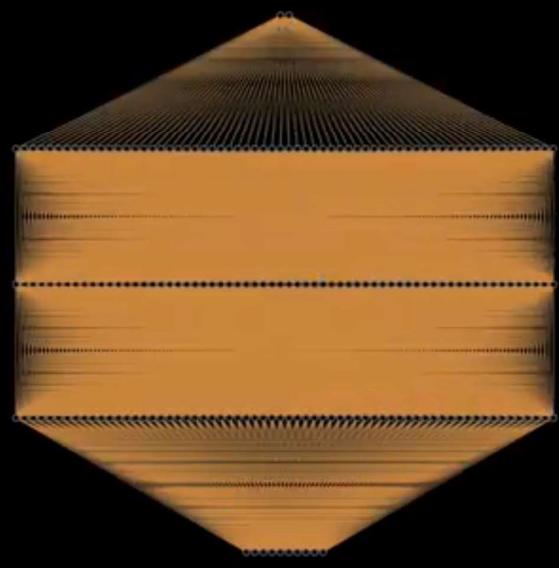
Layer sizes: 10, 32, 32, 32, 2



<https://nnfs.io>

is a huge number of uniquely tunable parameters

Layer sizes: 10, 64, 64, 64, 2



<https://nnfs.io>

9164 tunable parameters

```
1 import sys
2 import numpy as np
3 import matplotlib
4
5
6 print("Python:", sys.version)
7 print("NumPy:", np.__version__)
8 print("Matplotlib:", matplotlib.__version__)
```

```
Python: 3.7.7 (default, Mar 10 2020, 15:16:38)
```

```
[GCC 7.5.0]
NumPy: 1.18.2
Matplotlib: 3.2.1
[Finished in 0.3s]
```

```
1 inputs = [1.2, 5.1, 2.1]
2 weights = [3.1, 2.1, 8.7]
3 bias = 3
4
5
6
7
8 output = inputs[0]*weights[0] + inputs[1]*weights[1] + inputs[2]*weights[2] + bias
9
10 print(output)
```

```
1 inputs = [1, 2, 3]
2 weights = [0.2, 0.8, -0.5]
3 bias = 2
4
5
6
7
8
9 output = inputs[0]*weights[0] + inputs[1]*weights[1] + inputs[2]*weights[2] + bias
10 print(output)
```

```
2.3
[Finished in 0.0s]
```

```
inputs = [1.0, 2.0, 3.0]
weights = [0.2, 0.8, -0.5]
bias = 2.0

output = inputs[0]*weights[0] + inputs[1]*weights[1] + inputs[2]*weights[2] + bias
print(output)

>>> 2.3
```



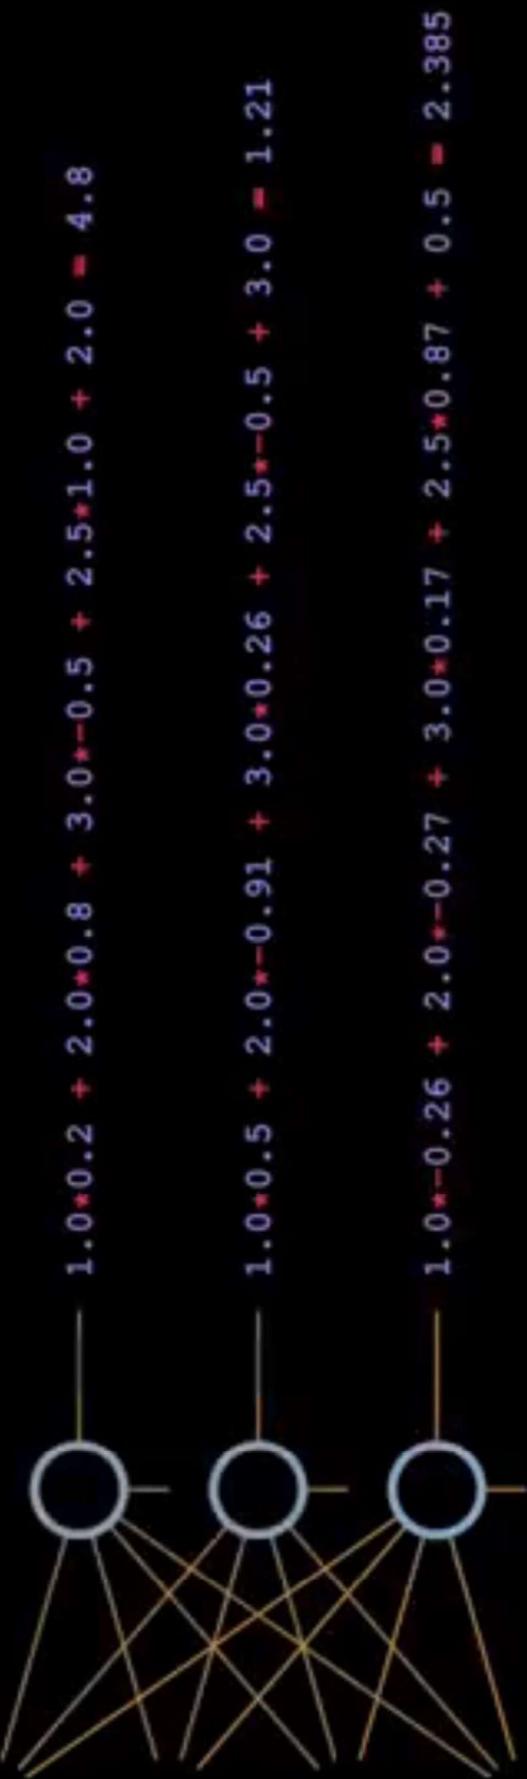
```

inputs = [1.0, 2.0, 3.0, 2.5]
weights1 = [0.2, 0.8, -0.5, 1.0]
weights2 = [0.5, -0.91, 0.26, -0.5]
weights3 = [-0.26, -0.27, 0.17, 0.87]
bias1 = 2.0
bias2 = 3.0
bias3 = 0.5

outputs = [
    inputs[0]*weights1[0] + inputs[1]*weights1[1] + inputs[2]*weights1[2] + inputs[3]*weights1[3] + bias1,
    inputs[0]*weights2[0] + inputs[1]*weights2[1] + inputs[2]*weights2[2] + inputs[3]*weights2[3] + bias2,
    inputs[0]*weights3[0] + inputs[1]*weights3[1] + inputs[2]*weights3[2] + inputs[3]*weights3[3] + bias3
]
print(outputs)

>>> [4.8, 1.21, 2.385]

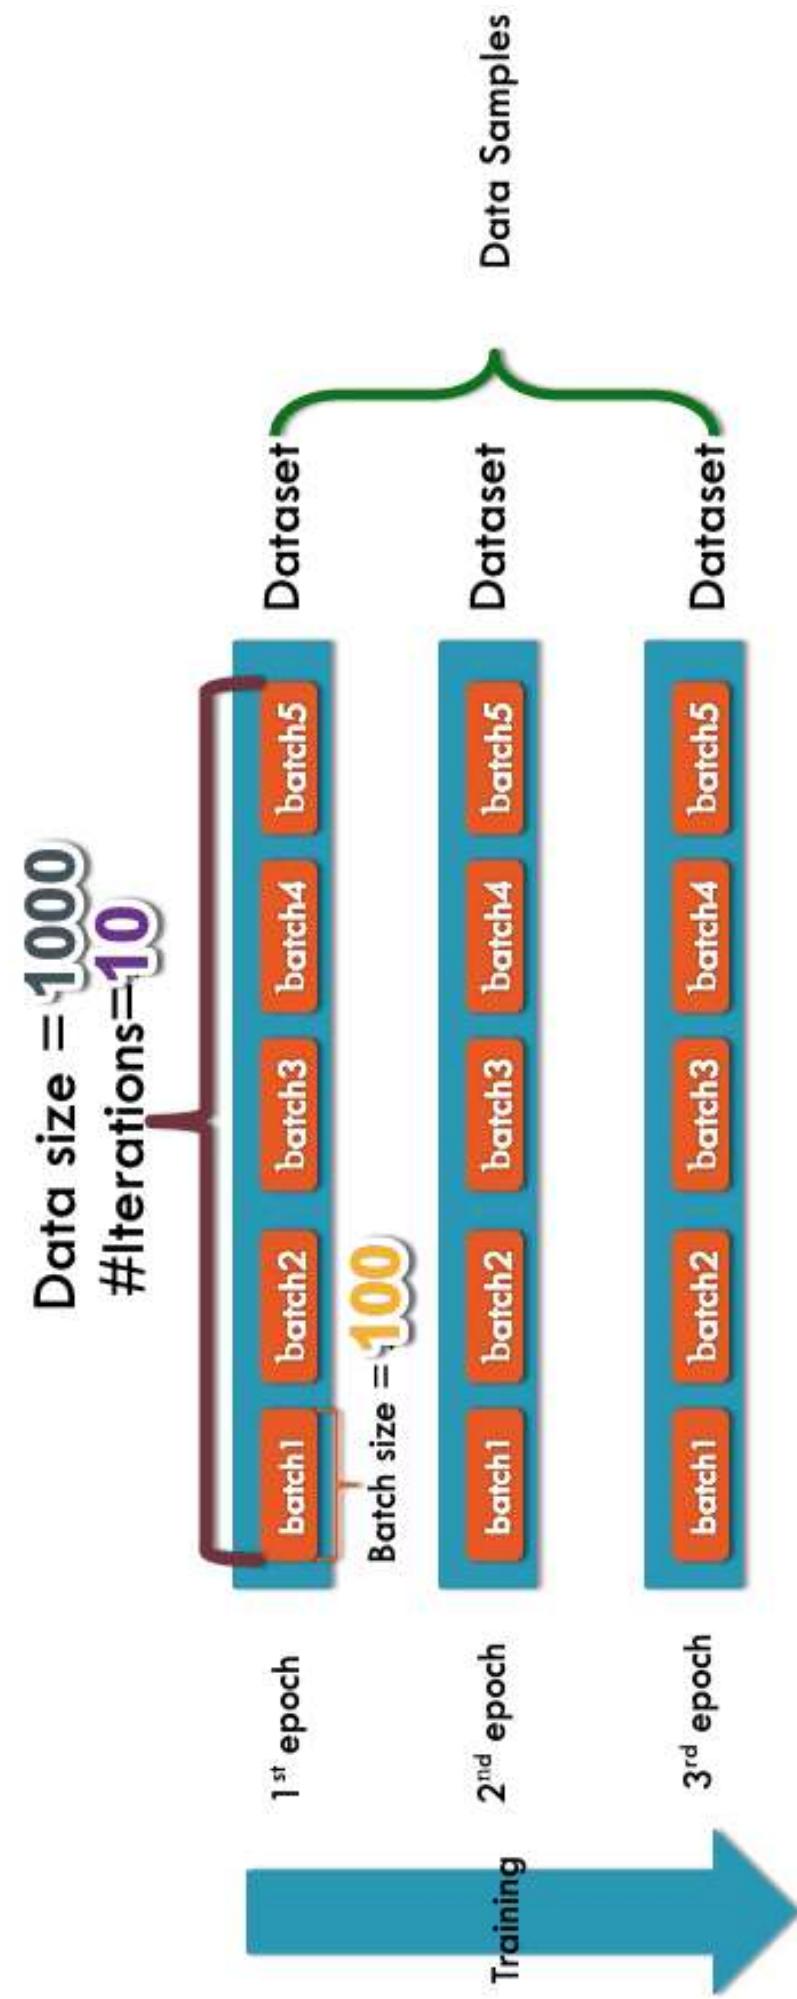
```



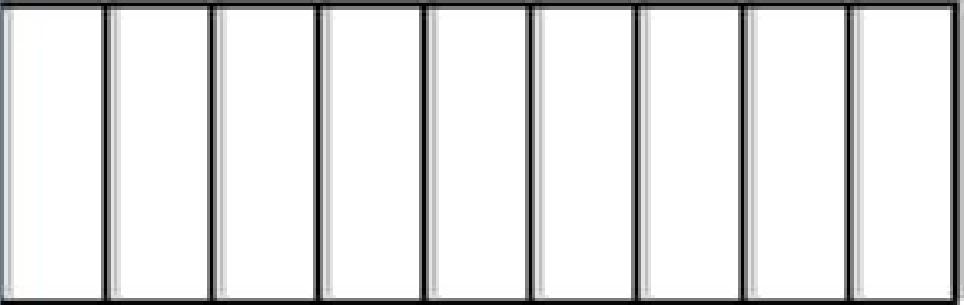
Epoch: One Epoch is when an ENTIRE dataset is passed forward and backward through the neural network only ONCE

Batch Size: Total number of training examples present in a single batch

Iteration: The number of passes to complete one epoch.

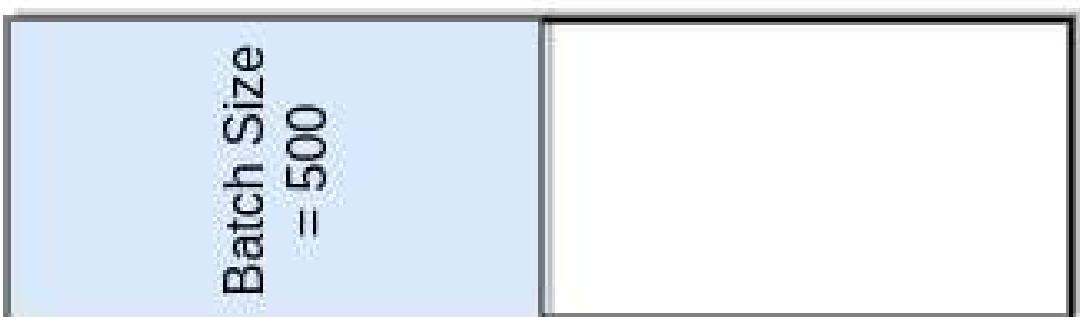


Batch Size =
100



Iterations per
Epoch = 10

Batch Size
= 500



Iterations per
Epoch = 2

Batch Size =
1000



Iterations per
Epoch = 1

Program 4: Build an artificial neural network by implementing the Backpropagation Algorithm and test the same using appropriate datasets.

import numpy as np

X = np.array(([2, 9], [1, 5], [3, 6]), dtype=float) *Normalizing the data*

y = np.array(([92], [86], [89]), dtype=float) *within data: 0 to 1*

X = X/npamax(X, axis=0) # maximum of X array longitudinally

y = y/100

#Sigmoid Function

def sigmoid (x):

return 1/(1 + np.exp(-x))

#Derivative of Sigmoid Function

def derivatives_sigmoid (x):

return x*(1-x)

#Variable initialization

```
epoch=7000 #Setting training iterations
```

```
lr=0.1 #Setting learning rate
```

```
inputlayer_neurons = 2 #number of features in data set
```

```
hiddenlayer_neurons = 3 #number of hidden layers neurons
```

```
output_neurons = 1 #number of neurons at output layer
```

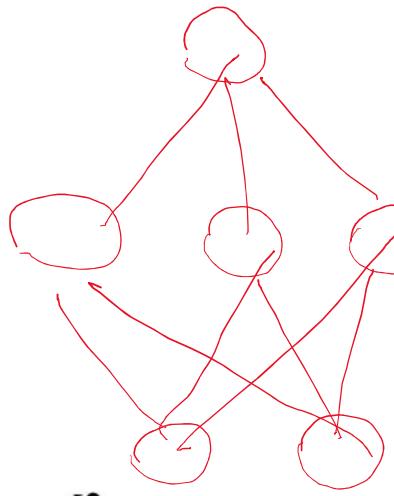
#weight and bias initialization

```
wh=np.random.uniform(size=(inputlayer_neurons,hiddenlayer_neurons)) QX3
```

```
bh=np.random.uniform(size=(1,hiddenlayer_neurons)) 1X3
```

```
wout=np.random.uniform(size=(hiddenlayer_neurons,output_neurons)) 1X1
```

```
bout=np.random.uniform(size=(1,output_neurons)) 1X1
```



~~3 times~~

```
#draws a random range of numbers uniformly of dim x*y
```

```
for i in range(epoch):
```

```
#Forward Propogation
```

```
hinp1=np.dot(X,wh)
```

```
hinp=hinp1 + bh
```

```
hlayer_act = sigmoid(hinp)
```

```
outinp1=np.dot(hlayer_act,wout)
```

```
outinp= outinp1+ bout
```

```
output = sigmoid(outinp)
```

#Backpropagation

```
EO = y-output
```

```
outgrad = derivatives_sigmoid(output)
```

```
d_output = EO* outgrad
```

```
EH = d_output.dot(wout.T)
```

```
hiddengrad = derivatives_sigmoid(hlayer_act)#how much hidden layer wts contributed to error
```

```
d_hiddenlayer = EH * hiddengrad
```

```
wout += hlayer_act.T.dot(d_output) * lr# dotproduct of nextlayererror and currentlayerop
```

```
wh += X.T.dot(d_hiddenlayer) * lr
```

```
print("Input: \n" + str(X))
```

```
print("Actual Output: \n" + str(y))
```

```
print("Predicted Output: \n" ,output)
```

//output

Input:

```
[[ 0.66666667 1.      ]  
 [ 0.33333333 0.55555556]  
 [ 1.      0.66666667]]
```

Actual Output:

```
[[ 0.92]  
 [ 0.86]  
 [ 0.89]]
```

Predicted Output:

```
[[ 0.90224607]  
 [ 0.87341151]  
 [ 0.8925683 ]]
```

4) Build an Artificial Neural Network by implementing the Backpropagation algorithm and test the same using appropriate data sets.

```

import numpy as np
X = np.array([[2, 9], [1, 5], [3, 6]], dtype='float')
y = np.array([[92], [86], [89]], dtype='float')
X = X/np.amax(X, axis=0) # maximum of
y = y/100

#Sigmoid Function
def sigmoid(x):
    return 1/(1 + np.exp(-x))

#Weight and bias initialization
wh=np.random.uniform(size=(inputlayer_neurons,hiddenlayer_neurons))
bh=np.random.uniform(size=(1,hiddenlayer_neurons))
wout=np.random.uniform(size=(hiddenlayer_neurons,output_neurons))
bout=np.random.uniform(size=(1,output_neurons))

```

#draws a random range of numbers uniformly of dim x*y Back propagation steps

for i in range(epoch):

#Forward Propogation

hinp1=np.dot(X,wh)

hinp=hinp1 + bh

hlayer_act = sigmoid(hinp)

outinp1=np.dot(hlayer_act,wout)

outinp=outinp1 + bout

output = sigmoid(outinp)

#Backpropagation

EO = y-output

outgrad = derivatives_sigmoid(output)

d_output = EO * outgrad

EH = d_output.dot(wout.T) # T : means transpose because here weight vector used

hiddengrad = derivatives_sigmoid(hlayer_act)#how much hidden layer wts contributed to error

d_hiddengrad = EH * hiddengrad

wout += hlayer_act.T.dot(d_output) * lr# dotproduct of nextlayererror and currentlayerop

bout += np.sum(d_output, axis=0,keepdims=True) * lr

wh += X.T.dot(d_hiddengrad) * lr

#bh += np.sum(d_hiddenlayer, axis=0,keepdims=True) * lr

print("Input: \n" + str(X))

print("Actual Output: \n" + str(y))

print("Predicted Output: \n" + str(wout))

$$\Delta w_{ji} = \eta \delta_j x_{ji}$$

Data Set : Play ⁱⁿ Study ⁱⁿ Sleep ⁱⁿ
 Features (Hrs Slept, Hrs Studied)
 Labels (marks obtained)

Build an Artificial Neural Network by implementing the Backpropagation algorithm and test the same using appropriate data sets.

4) Build an Artificial Neural Network by implementing the Backpropagation algorithm and test the same using appropriate data sets.

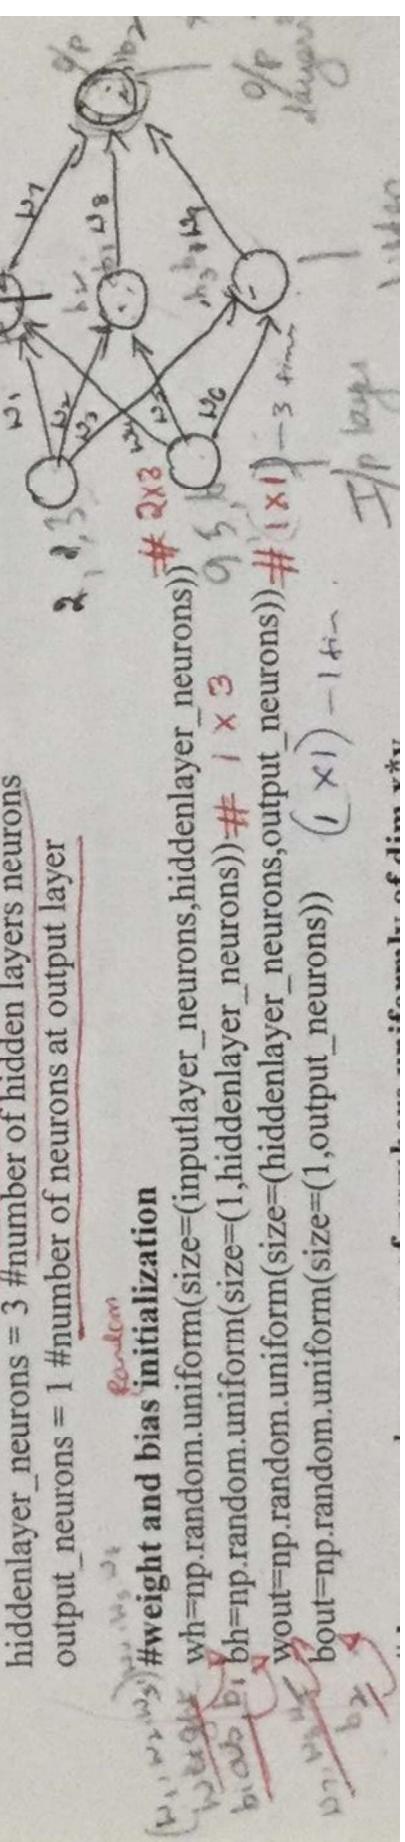
```

import numpy as np
X = np.array([[2, 9], [1, 5], [3, 6]], dtype=float)
y = np.array([[92], [86], [89]], dtype=float)
X = X/np.amax(X, axis=0) # maximum of X array longitudinally
y = y/100 # normalise output between 0 to 1.
#Sigmoid Function
def sigmoid(x):
    return 1/(1 + np.exp(-x))

```

```
#Sigmoid Function
def derivatives_sigmoid(x):
    return x*(1-x) - (x * (1 - x)) in one]
```

```
#Variable initialization
epoch=7000 #Setting training iterations
lr=0.1 #Setting learning rate
#answer: not "the net"
inputlayer_neurons = 2 #number of features in data set
hiddenlayer_neurons = 3 #number of hidden layers neurons
output_neurons = 1 #number of neurons at output layer
```



#draws a random range of numbers uniformly or dim xy

for i in range(epoch):

#Forward Propagation

himp1 = np.dot(X, wh)
himp = himp1 + bh
hlayer_act = sigmoid(himp)
outinp1 = np.dot(hlayer_act, wout)
outinp = outinp1 + bout
output = sigmoid(outinp)

#Backpropagation

EO = y - output
outgrad = derivatives sigmoid(output)
d_output = EO * outgrad
EH = d_output.dot(wout.T)
hiddengrad = derivatives sigmoid(hlayer_act)
d_hiddenlayer = EH * hiddengrad
wout += hlayer_act.T.dot(d_output)
bout += np.sum(d_output, axis=0, keepdims=True) * lr
wh += X.T.dot(d_hiddenlayer) * lr
#bh += np.sum(d_hiddenlayer, axis=0, keepdims=True) * lr
print("Input: \n" + str(X))
print("Actual Output: \n" + str(y))

#dot product + bias layer
himp = np.dot(x, wh) + bh
hlayer_act = sigmoid(himp)
outinp1 = np.dot(hlayer_act, wout)
outinp = outinp1 + bout
output = sigmoid(outinp)
Backpropagation
EO = y - output # Error at output layer
outgrad = derivatives sigmoid(output) # derivative of sigmoid
d_output = EO * outgrad # derivative of output
EH = d_output.dot(wout.T) # means transpose
hiddengrad = derivatives sigmoid(hlayer_act) # how much hidden layer wts contributed to error
d_hiddenlayer = EH * hiddengrad # how much hidden layer wts contributed to error
wout += hlayer_act.T.dot(d_output) # dotproduct of nextlayererror and currentlayerop
bout += np.sum(d_output, axis=0, keepdims=True) * lr
wh += X.T.dot(d_hiddenlayer) * lr
#bh += np.sum(d_hiddenlayer, axis=0, keepdims=True) * lr
print("Input: \n" + str(X))
print("Actual Output: \n" + str(y))
transpose of want
dot product with d-output.
Take transpose first & then dot product with d-output.
hlayer_act.T. dot (d_output) \Rightarrow X.T. d-hiddenlayer
X.T. dot (d_hiddenlayer) \Rightarrow X.T. d-hiddenlayer

```
print("Predicted Output: \n",output)
```

//output

Input:

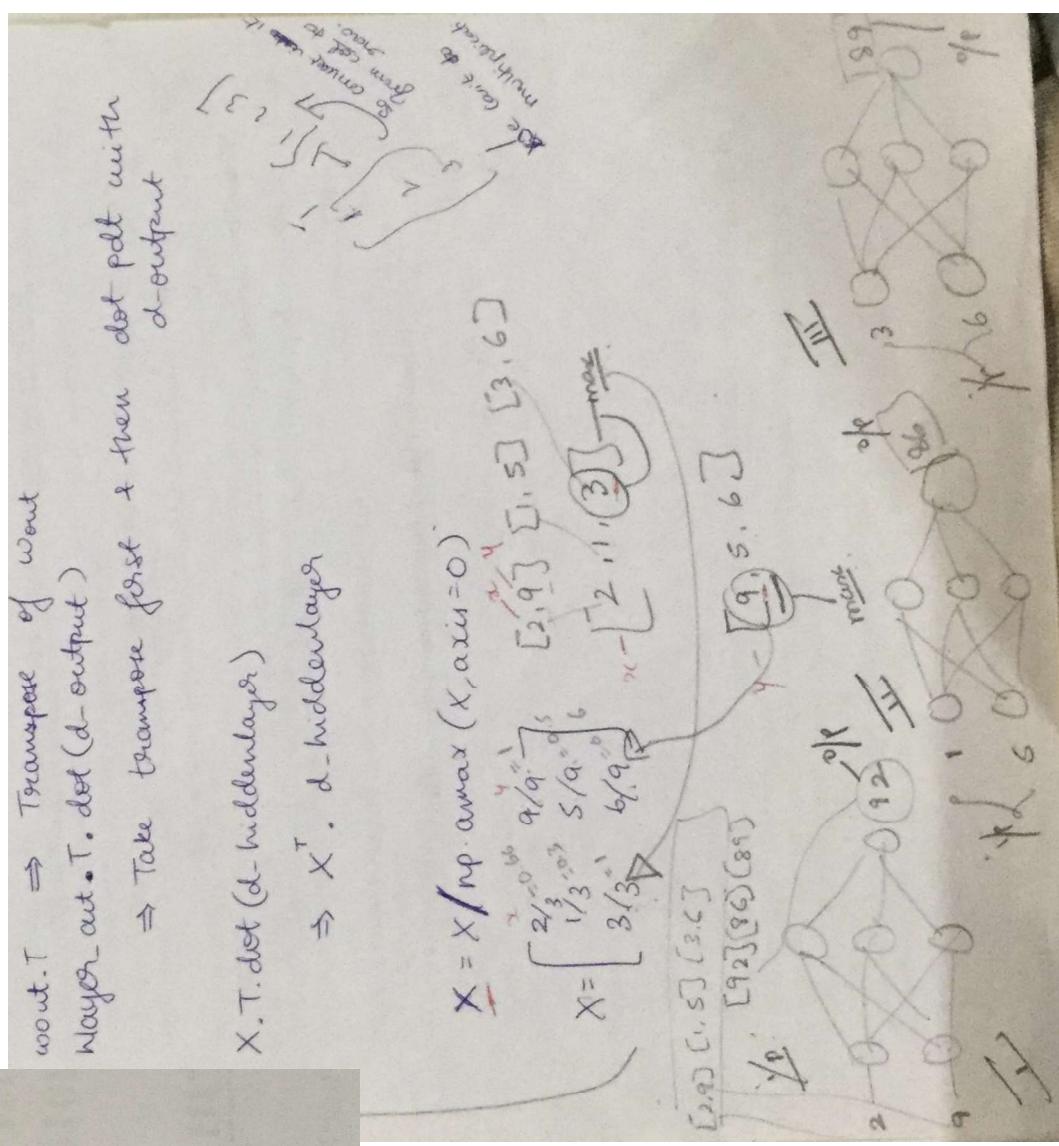
```
[[ 0.66666667  1.
   [ 0.33333333  0.55555556]
   [ 1.          0.66666667]]]
```

Actual Output:

```
[[ 0.92]
 [ 0.86]
 [ 0.89]]
```

Predicted Output:

```
[[ 0.90224607]
 [ 0.87341151]
 [ 0.8925683 ]]
```



4) Build an Artificial Neural Network by implementing the Backpropagation algorithm and test the same using appropriate data sets.

```

import numpy as np
X = np.array([[2, 9], [1, 5], [3, 6]], dtype=float)
y = np.array([[92], [86], [89]], dtype=float)
X = X/np.amax(X, axis=0) # maximum of X array longitudinally
y = y/100

#Sigmoid Function
def sigmoid(x):
    return 1/(1 + np.exp(-x))

#Sigmoid Function // derivative of Sigmoid func.
def derivatives_sigmoid(x):
    return x*(1-x)

```

random uniform(a, b)
 \downarrow
 generate random floating \mathcal{Z} .

Print $\frac{N}{\text{no. of } \mathcal{Z}}$
 such that (\mathcal{Z}^1)
 $a \leq N \leq b$ if $a \leq b$
 $b \leq N \leq a$ if $b \leq a$

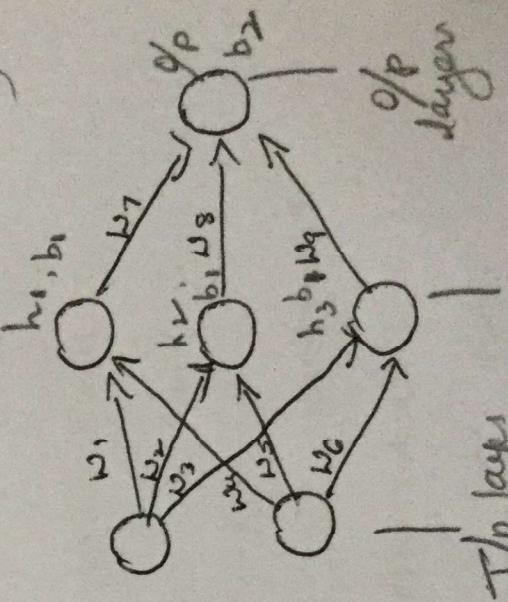
#Variable initialization
 epoch=7000 #Setting training iterations
 lr=0.1 #Setting learning rate
 inputlayer_neurons = 2 #number of features in data set
 hiddenlayer_neurons = 3 #number of hidden layers neurons
 output_neurons = 1 #number of neurons at output layer

#weight and bias initialization

```

wh=np.random.uniform(size=(inputlayer_neurons,hiddenlayer_neurons))
bh=np.random.uniform(size=(1,hiddenlayer_neurons))
wout=np.random.uniform(size=(hiddenlayer_neurons,output_neurons))
bout=np.random.uniform(size=(1,output_neurons))

```



#draws a random range of numbers uniformly of dim x*y

for i in range(epoch):

#Forward Propogation
 ~~if weights~~ $\rightarrow \{w_i * x_i\}$
 ~~weights~~ $\rightarrow h_{inp} = np.dot(X, wh)$
 ~~bias~~ $\rightarrow h_{inp} + bh$
 ~~hidden layer act~~ $\rightarrow h_{layer_act} = sigmoid(h_{inp})$
 ~~outinp1~~ $\rightarrow outinp1 = np.dot(h_{layer_act}, wout)$
 ~~outinp1 + bout~~ $\rightarrow outinp1 + bout$
 ~~outinp~~ $\rightarrow outinp = sigmoid(outinp1 + bout)$
 ~~output~~ $\rightarrow output = sigmoid(outinp)$

#Backpropagation
 ~~predicted value~~ $\rightarrow EO = y - output$
 ~~computed target~~ $\rightarrow EO$
 ~~outgrad~~ $\rightarrow outgrad = derivatives_sigmoid(output)$
 ~~d_output~~ $\rightarrow d_output = EO * outgrad$
 ~~EH~~ $\rightarrow EH = d_output.dot(wout.T)$
 ~~hiddengrad~~ $\rightarrow hiddengrad = derivatives_sigmoid(hlayer_act)$
 ~~hiddengrad~~ $\rightarrow d_hiddenlayer = EH * hiddengrad$
 ~~wout~~ $\rightarrow wout += hlayer_act.T.dot(d_output) * lr$
 ~~# bout~~ $\rightarrow # bout += np.sum(d_output, axis=0, keepdims=True) * lr$
 ~~wh~~ $\rightarrow wh += X.T.dot(d_hiddenlayer) * lr$
 ~~#bh~~ $\rightarrow #bh += np.sum(d_hiddenlayer, axis=0, keepdims=True) * lr$
 print("Input: \n" + str(X))
 print("Actual Output: \n" + str(y))

```
print("Predicted Output: \n",output)
```

```
//output
```

```
Input:
```

```
[ [ 0 . 666666667 1 .  
[ 0 . 333333333 0 . 55555556 ]  
[ 1 . 0 . 666666667 ] ]
```

```
Actual Output:
```

```
[ [ 0 . 92 ]  
[ 0 . 86 ]  
[ 0 . 89 ] ]
```

```
Predicted Output:
```

```
[ [ 0 . 90224607 ]  
[ 0 . 87341151 ]  
[ 0 . 8925683 ] ]
```